



UNIVERSITY OF
CAMBRIDGE

Department of Computer
Science and Technology

Accelerating Molecular Generation through Representation Alignment

Shreyas Ravishankar

Hughes Hall

June 2026

Submitted in partial fulfillment of the requirements for the
Master of Philosophy in Advanced Computer Science

Total page count: 60

Main chapters (excluding front-matter, references and appendix): 45 pages (pp 7–51)

Main chapters word count: 14343

Methodology used to generate that word count:

The count is produced with `texcount` over the $\text{\LaTeX}$ source, following `\input` files. It is scoped to the core chapters by `%TC:ignore` directives that exclude the front matter, bibliography, and appendix. The reported figure sums body text, section headings, and figure/table captions (i.e. narrative text); tabular data listings are not counted, in keeping with the word-count rules. Reproduce with:

```
$ make wordcount
perl tools/texcount.pl -inc -sum -total -1 report-draft.tex
14343
```

Declaration

I, Shreyas Ravishankar of Hughes Hall, being a candidate for the Master of Philosophy in Advanced Computer Science, hereby declare that this report and the work described in it are my own work, unaided except as may be specified below, and that the report does not contain material that has already been used to any substantial extent for a comparable purpose. In preparation of this report, I adhered to the Department of Computer Science and Technology AI Policy. [The project required the approved use of {insert name of technology} and such use is acknowledged in the text at the relevant sections.] [I am content for my report to be made available to the students and staff of the University.]

Signed:

Date:

Abstract

Write a summary of the whole thing. Make sure it fits on one page.

Contents

1	Introduction	7
1.1	A history of trade-offs	7
1.2	Representation alignment	8
1.3	Contributions	9
2	Background and Related Work	10
2.1	Flow matching	10
2.2	Representation alignment	11
2.2.1	What makes for a good REPA encoder?	13
2.3	A primer on molecular generation	13
2.4	Tabasco and Proteina	14
2.4.1	Small molecules and Tabasco	15
2.4.2	Protein backbones and Proteina	15
2.5	Open questions	16
3	Evaluating molecular generation	17
3.1	What makes for a “good” molecule?	17
3.2	Representation and generation are multi-factorial	17
3.3	The linear probe principle	18
3.4	Small-molecule metrics	19
3.5	Protein-backbone metrics	19
3.6	Intepreting trade-offs to evaluate models	23
4	Encoder targets and profiling	24
4.1	A profiling protocol	24
4.2	Small molecule encoders	26
4.3	Protein backbone encoders	27
4.4	Informing our experiments	28
5	REPA on small molecules	29
5.1	Experimental setup	29
5.1.1	Dataset	29
5.1.2	Models	29

5.1.3	Metrics	30
5.2	Results: no measurable gain	31
5.3	Discussion	32
6	REPA on protein backbones	35
6.1	Experimental setup	35
6.1.1	Datasets	35
6.1.2	Models	37
6.1.3	Metrics	37
6.1.4	Evaluation protocol	39
6.2	Results	40
6.2.1	REPA improves representation quality asymptotically	40
6.2.2	A diagnostic check on alignment	41
6.2.3	REPA accelerates generation quality by climbing the representation– generation envelope faster	42
6.2.4	REPA improves whole-set coverage but may trade off diversity within its designable subset	45
6.2.5	REPA’s gains hold across sampler noise levels	47
6.2.6	What we do not claim	48
6.2.7	Robustness across scale	49
6.3	Discussion	49
7	Conclusions	51
A	Technical details	57
A.1	Dataset composition and split overlap	57
A.2	Choice of CATH fold-probe level	57
A.3	Leakage-controlled probe evaluation	58
A.4	Representation-alignment matrix	59
A.5	Representation–generation coupling on AFDB	59
A.6	The ssJSD-2D secondary-structure metric	59
A.7	Secondary-structure composition of the designable subset	60
A.8	Full encoder profiles	60

Chapter 1

Introduction

The evolution of generative machine learning in scientific domains is defined by the shifting tensions between data scale, data quality, and structural priors. Our report examines a recent addition to this story, *representation alignment*, in two regimes for de novo molecular generation — small molecules and protein backbones — and characterises its behaviour. In this introduction, we trace the narrative arc of research trends in molecular generation, and highlight the motivation for our work: the emerging need for data-efficient modelling in a world of prior-light models.

1.1 A history of trade-offs

The dominant ideas in generative modelling are *denoising diffusion* and its offshoot *flow matching* [1, 2], which found breakthrough success in the realm of image generation. This paradigm has since been adapted for de novo design challenges across scientific fields, such as generating protein backbones [18, 19, 14], inorganic materials [21], and drug-like small molecules [15]. Training these models is typically expensive in both data and compute, with state-of-the-art image generators routinely training on billions of examples [34]. However, the corpus of available high-quality molecular data sits orders of magnitude below this. Experimental determination is slow and expensive, which means datasets are correspondingly small. For example, protein structure generators have largely been trained on at most half a million structures [14], which may explain why they have not yet matched the maturity of their vision-domain counterparts.

Historically, researchers compensated for the small scale of high-quality molecular data by encoding structural priors directly into model architectures [40]. SE(3)-equivariant architectures, for example, enforce the rotational and translational symmetries expected of molecules by construction [42]. But hard structural priors come with a computational cost, and the operations required to encode geometry are a bottleneck to scaling models. A recent wave of research has moved entirely the other way: drop hard priors, use bare-bones non-equivariant transformer stacks, and rely on data scale and augmentation to teach the model what equivariance buys. This shift fits a broader trend in machine

learning on leveraging representation instead of model architecture to inject soft inductive biases for structured data. Vision Transformers [43] showed that the spatial-locality and translation-equivariance priors of convolutional neural nets (CNNs) can be learned rather than imposed, given sufficient scale. Graph Transformers like TokenGT [44] have made the same case for graphs. AlphaFold 3 [46] stripped many of the SE(3)-equivariant biases of AlphaFold 2 [45], replacing the structure module with a diffusion model operating on raw coordinates. Brehmer et al. [47] make the empirical scaling case directly: more data through a simpler model beats less data through a richer one. In our work, we examine Tabasco [15] and Proteina [14], which are flow-based generators built on largely vanilla transformers for small molecules and protein backbones respectively.

Given the difficulty of procuring data in scientific domains, data scale must necessarily come from non-experimental sources, which may be of lower quality. For example, recent efforts in protein generation have leaned on synthetic structures from folding-model predictions [14, 36] to expand training datasets into the tens of millions. However, even these enhanced corpora are arguably insufficient. Metagenomic databases catalogue over two billion natural protein sequences [38], and the theoretical sequence space at a median protein length [39] is hundreds of orders of magnitude greater still. Even at the frontier of the field’s data scale, we sample but a sliver of the distribution we wish to model.

In summary, the field has traded hard structural priors for architectural simplicity, and curated data for synthetic scale. However, this trade-off dilutes the structural signal available to a model. Furthermore, even our best attempts at creating large datasets may not be large enough. These twin challenges amplify the need for techniques that make training still more efficient without sacrificing quality.

1.2 Representation alignment

But why is training diffusion-transformer architectures expensive in the first place? One answer comes from Yu et al. [5], who argue that these models suffer from a *representation bottleneck*. During training, a model is implicitly asked to learn two distinct tasks simultaneously: (1) representation: extracting semantically meaningful encodings from complex data; and (2) generation: constructing data from these encodings. The standard denoising objective alone may not provide sufficient signal for representation learning, leading to slow convergence and the large training budgets these models require in practice.

This burden falls harder on the prior-light models we study. In the absence of equivariant architectures, the training objective must now shoulder the additional responsibility of learning structural symmetries. Moreover, the scarcity of high-quality data compounds the problem. The model must learn complex geometry from data that only sparsely samples the true distribution, with no architectural scaffolding to fall back on.

One response to this problem is to supply representational structure from elsewhere. Instead of being entirely left to the training objective, it can be injected softly as an

auxiliary training signal. This is the idea behind Representation Alignment (REPA), which Yu et al. propose as a remedy. The technique adds a regularisation term that aligns a generator’s intermediate hidden states with the embeddings of a frozen pretrained encoder. On image diffusion, they report roughly a $17.5\times$ training speed-up against a strong baseline to achieve similar generation quality. REPA has since been extended beyond images to video [9], audio [10], and recently to molecular generation [7, 8], where it has been demonstrated but not systematically characterised.

1.3 Contributions

In this light, the core question we investigate in this report is: *can prior-light molecular generation models benefit from representation alignment with pre-trained models, in training efficiency and generation quality?* We aim to characterise when and why this works, including failure modes and limitations.

We document two lines of enquiry.

- (i) *Does REPA work in the molecular domain?*: We conduct empirical studies on training flow-based generators with REPA in two regimes: Tabasco [15] for small molecules, and Proteina [14] for protein backbones. We find that REPA does not measurably improve Tabasco, but it improves training efficiency and advances the generation quality envelope for Proteina.
- (ii) *When does REPA work?*: Following Singh et al. [6] in the realm of image generation, we develop a profiling framework to characterise the suitability of an encoder as a representation target. We measure three properties of an encoder: (1) an upper bound on the structural information it contains (linear-probe recoverability against domain structural labels); (2) a lower bound on what a projector can transfer from that information (projector-based recoverability); (3) the geometric conditioning of its embeddings for the alignment objective (effective rank).

We do not propose a new architecture or alignment loss, and our two studies are not a controlled experiment over the conditions that govern REPA’s effectiveness. Instead, they are best read as paired case studies in regimes where those conditions differ.

The remainder of this report is organised as follows. Chapter 2 sets out the background and context needed to read the rest: flow matching, REPA, molecular generation. Chapter 3 discusses what it means for a generative model to be “good” in our domain, and the tensions that question hides. Chapter 4 develops the encoder-profiling framework and uses it to predict where REPA may and may not help. Chapters 5 and 6 present the two paired studies, small molecules first and protein backbones second. Finally, we conclude in Chapter 7 with a discussion of whether what we learn generalises beyond molecular generation.

Chapter 2

Background and Related Work

In this chapter, we introduce the machinery the rest of this thesis depends on. We outline flow matching (§2.1), the generative paradigm underlying the models we study, and representation alignment (§2.2), the regulariser that is our object of study. We then turn to the molecular setting. We situate generative models within the broader de novo molecular design pipeline (§2.3), before presenting the two molecular domains we work in, small molecules and protein backbones, and the flow-matching models we extend, Tabasco and Proteina (§2.4). The treatment is cursory by design. We intend only to guide the reader through this report, and we provide citations for omitted details.

2.1 Flow matching

At first glance, the term *generative modelling* implies the problem of *creating* new artefacts from scratch. However, it is perhaps more accurately described as the problem of *sampling* new artefacts from a complex data distribution that can only be accessed through a finite training set. A common strategy here is to fix a simple distribution we can sample from cheaply — typically standard Gaussian noise $p_0 = \mathcal{N}(0, I)$ — and learn a mapping from it to the desired data distribution p_1 .

Flow matching (FM) [2] realises this mapping as a continuous-time transport. The model designer defines a *probability path* between a noise sample and a data sample — a continuous interpolation between them — and the model learns to carry the noise sample along this path. Concretely, it parameterises a time-dependent *velocity field* $v_\theta(x, t)$ that gives the instantaneous direction of motion at each location x and time $t \in [0, 1]$. At inference, integrating the learned ordinary differential equation (ODE) $\dot{x}_t = v_\theta(x_t, t)$ from $t = 0$ to $t = 1$ pushes a fresh noise sample forward into a fresh data sample.

A common choice of probability path is the linear interpolant, $x_t = (1 - t)x_0 + tx_1$ for $x_0 \sim p_0$ and $x_1 \sim p_1$. The model’s target is the *marginal* velocity field at (x, t) : the average direction of motion over all noise-data pairs whose paths pass through that point. However, supervising this marginal directly is intractable, since the contributing pairs are not known in advance. Flow matching sidesteps this difficulty with a key observation [2]:

although the marginal field may be intractable, the *conditional* velocity along the path of any specific noise-data pair is relatively trivial. Indeed, for the linear interpolant, it is simply $\dot{x}_t = x_1 - x_0$. Regressing the model on these conditional velocities, averaged over many sampled pairs, recovers the marginal field in expectation, and gives the training loss as a mean-squared regression:

$$\mathcal{L}_{\text{FM}} = \mathbb{E}_{t, x_0, x_1} \|v_\theta(x_t, t) - (x_1 - x_0)\|^2.$$

In practice, every training step samples a time $t \in [0, 1]$ and a noise-data pair (x_0, x_1) , computes the interpolant x_t , and takes a gradient step on this squared error.

Flow matching with this linear interpolant generalises score-based diffusion: the variance-preserving Gaussian-path objective of denoising diffusion models [1] is a special case of the FM family [2, 4].

Generating samples from a trained flow-matching model amounts to integrating the learned velocity field. This is often performed with a standard ODE solver, with Euler steps usually sufficing for stable trajectories. The field also admits stochastic sampling variants — such as those drawing on the connection to score-based stochastic differential equations (SDEs) — which inject noise during the integration steps to explore the probability space and improve sample diversity. The amount of injected noise is itself a tunable knob: more noise broadens diversity at some cost to per-sample quality, with the deterministic ODE recovered as its zero-noise limit. The models we work with in this report, Tabasco and Proteina, support both deterministic and stochastic samplers.

Conditional generation augments v_θ with context tokens — such as a class label, a sequence length, or a structural motif — with conditioning supplied at both training and sampling time. In this report, we focus on unconditional generation only.

For a more complete treatment of flow matching, we refer the reader to Lipman et al. [2] for the original presentation, and to the recent Lipman et al. guide [4].

2.2 Representation alignment

Chapter 1 introduced the argument made by Yu et al. [5] about the *representation bottleneck* in diffusion-transformer models. Their analysis applies to flow-matching models as well, given the equivalence between the two objectives discussed above (§2.1). The flow-matching objective implicitly asks a model to learn two things at once: a useful representation of the data, and a velocity field that interprets this representation. This twin burden, they argue, leads to slow convergence and large training budgets.

Their proposed solution, Representation Alignment (REPA), addresses the first half of this burden by adding a regularisation term to the training loss. Figure 2.1 shows the construction. The regularisation pulls z_ℓ , the generator’s intermediate hidden state at a chosen *alignment layer* ℓ , towards the embeddings produced by a frozen pretrained




Figure 2.1: The REPA construction. A noised sample $x_t = (1-t)x_0 + tx_1$ is constructed at each training step from a noise sample x_0 , a clean data sample x_1 , and a sampled timestep $t \in [0, 1]$. The generator v_θ processes x_t through its sequence of layers $\ell_1, \dots, \ell_L$ to produce the velocity field that feeds the flow-matching loss $\mathcal{L}_{\text{FM}}$. Separately, the clean sample x_1 is also passed through a frozen pretrained encoder f_{enc} (shaded). The generator’s hidden state z_ℓ at the chosen alignment layer is mapped into the encoder’s feature space by a trainable projector h_ϕ ; the REPA loss $\mathcal{L}_{\text{REPA}}$ is the negated similarity between $h_\phi(z_\ell)$ and $f_{\text{enc}}(x_1)$, averaged across tokens. The total training loss $\mathcal{L}$ combines them with weight λ .

encoder f_{enc} run on the clean data sample x_1 . The hidden state z_ℓ is a sequence of N token vectors. In the context of images, each token represents an image patch; in the molecular setting, each token represents an atom or residue that is part of the larger molecule (§2.3). A small trainable *projector head* h_ϕ — a three-layer multi-layer perceptron (MLP) in the original paper — maps z_ℓ into the encoder’s feature space, so that $h_\phi(z_\ell)$ and $f_{\text{enc}}(x_1)$ can be compared. This alignment is asymmetric. The encoder provides a fixed target throughout training, while the generator parameters θ and the projector ϕ are updated.

The alignment loss is the negated similarity, averaged token-wise, between projected and target embeddings:

$$\mathcal{L}_{\text{REPA}} = -\mathbb{E}_{t, x_0, x_1} \left[\frac{1}{N} \sum_{n=1}^N \text{sim}(h_\phi(z_\ell)_n, f_{\text{enc}}(x_1)_n) \right],$$

where n indexes the N tokens within a sample. In the image setting, N is constant across the dataset, so this token-wise average is equivalent to averaging similarities once per sample and then across the batch. Indeed, the original paper draws no distinction between the two. For molecular data, where N varies per sample, the two diverge. We return to this distinction in the design knobs below.

The final training objective is the flow-matching loss from §2.1 plus this term, scaled by a weight λ :

$$\mathcal{L} = \mathcal{L}_{\text{FM}} + \lambda \mathcal{L}_{\text{REPA}}.$$

This construction exposes several design choices, some of whose effects we revisit in the ablations of Chapter 6:

- The choice of encoder f_{enc} . Different encoders may capture distinct aspects of the data distribution.
- The alignment layer ℓ at which we attach the projector. Yu et al. experimentally determined that mid-network layers work best for image diffusion, but the optimal choice is not obvious a priori.
- The alignment-loss weight λ , which balances the alignment signal against the generative objective. Once again, Yu et al. experimentally determined an optimal value of 0.5 for image diffusion, but this may not hold universally.
- The choice of similarity metric sim . Yu et al. compared cosine similarity and mean-squared error for image diffusion, and found the former more effective. We adopt cosine similarity in this report as well.
- The averaging mode. Averaging tokens within each sample first weights every molecule equally, while pooling tokens across the batch gives larger molecules proportionally more REPA signal.

We adopt the core REPA construction as in Yu et al. for this report. However, a small family of methodological variants on REPA has since emerged. Dispersive regularisers [12] replace the external encoder with an internal contrastive term, while flexible-target formulations [7] broaden which encoder features the loss attaches to. Separately, REPA-style alignment has been applied beyond the image setting, to video [9], audio [10], materials [11], and molecules [8]. We defer exploration of these variants to future work.

2.2.1 What makes for a good REPA encoder?

There are many possible choices for the encoder f_{enc} in any REPA construction, and this selection is a key design decision. But what characteristics make for a good alignment target? In their original report, Yu et al. [5] observed that ImageNet linear-probe accuracy correlates positively with generation performance, suggesting that global semantic representation is the driving factor. However, more recent work by Singh et al. [6] reports the opposite after profiling a wider range of encoders. Linear-probe accuracy does not in fact predict REPA gain, they argue, while a measure of the spatial structure of patch-token representations does. This suggests that generation performance for REPA may be driven by *local* feature geometry, rather than *global* semantic strength. Encoder characterisation for REPA remains an open research question, particularly in the molecular generation setting, and we pick up the thread in Chapter 4.

2.3 A primer on molecular generation

Having outlined the modelling paradigms driving our report, we now turn to the molecular setting in which we apply them. To situate our work, we begin with a brief primer on




Figure 2.2: The de novo molecular design cycle. An upstream specification drives a generative model to produce a candidate pool, which is funnelled through in-silico filters and then synthesised (or expressed) and assayed in the wet lab. Experimental results feed back into the next round of generation. Typical generation-step pool sizes are in the thousands to tens of thousands of candidates per round [18].

the role generative models play within the broader de novo molecular design pipeline. Generative models are rarely intended to produce ready-to-use artefacts. Instead, they act as high-throughput in-silico proposal engines in an iterative process of refinement and selection. This process, often called a *design campaign* [55], is summarised in Figure 2.2.

An upstream specification, such as a drug-likeness profile or a fold class, conditions a generative model to produce a pool of candidate samples. These pools typically comprise thousands to tens of thousands of candidates per round [18], which are then funnelled through in-silico filters. For small molecules, these typically include checks for structural validity and physical plausibility [56], drug-likeness [59], and target docking; for protein backbones, designability and structural diversity are the dominant considerations [14]. The select few survivors are produced in a laboratory, through chemical synthesis (small molecules) or expression in cell lines (proteins), and evaluated using in vitro or in vivo assays. Experimental results then feed back into the next round of generation.

This pipeline framing is relevant to our discussion of evaluation metrics in Chapter 3. Several metrics we adopt, such as structural validity and designability, are not measures of intrinsic molecular quality, but proxies for specific in-silico filter steps.

2.4 Tabasco and Proteina

The two models we study — Tabasco [15] for small molecules and Proteina [14] for protein backbones — are both non-equivariant transformer flow-matching generators. They are explicitly designed to scale with data and model size, rather than relying on structural priors imposed by model architecture. The rotational and translational symmetries expected of 3D molecules are intended to be learned from data and augmentation rather than imposed by construction.

Tabasco is motivated by the success of similar architectures in protein-structure generation, like Proteina, setting it apart from equivariant message-passing small-molecule generators such as EQGAT-diff [50], MiDi [51], and SemlaFlow [52]. In turn, Proteina

is motivated by the pair-bias attention mechanism introduced in AlphaFold 3 [46], which distinguishes it from SE(3)-equivariant backbone generators such as RFDiffusion [18], FrameFlow [48], and Genie [49]. Despite their architectural minimalism, both achieve performance comparable or superior to equivariant baselines on structural-quality evaluations.

2.4.1 Small molecules and Tabasco

In medicinal chemistry, a *small molecule* is a drug-like compound at the sub-900-Dalton scale [53]. *in-silico*, its structure is typically represented using two coupled features: a discrete *molecular graph* of typed atoms connected by typed bonds, and a continuous *conformer* placing those atoms in 3D Cartesian space. A flow-matching parameterisation must therefore handle both jointly.

Tabasco [15] represents each molecule as a sequence of per-atom tokens, each carrying an atom-type embedding and a 3D coordinate, but with no explicit bond decoder. This treats the molecule as a point cloud of typed atoms rather than a graph, sidestepping edge modelling at the cost of having to recover bond structure post hoc from interatomic distances. Tabasco is trained on GEOM-drugs [37], a public dataset of drug-like conformers.

This point-cloud-of-atoms parameterisation is one of several options in the literature, and each makes a different trade-off. *SMILES* and *SELFIES* are string encodings that work well with autoregressive models but discard the 3D conformer entirely; we encounter them later as the parameterisation employed by some candidate encoders we evaluate in Chapter 4. *3D voxel grids* preserve geometry, but lose the discrete graph and inflate the token count by an order of magnitude or more. *Equivariant point clouds*, used by generators such as EQGAT-diff [50], MiDi [51], and SemlaFlow [52], keep both graph and geometry, but pay the computational cost associated with SE(3)-equivariant layers. We refer the reader to Duval et al. [41] and Wang et al. [22] for broader surveys in this space.

2.4.2 Protein backbones and Proteina

A *protein* is a chain of amino-acid residues, ranging from around 50 to several thousands in length [54]; the largest known, titin, has approximately 34,000 residues. (Shorter chains are typically classified as *peptides* instead.) Each amino-acid residue contributes four *backbone* atoms (N, C $_{\alpha}$, C, O) whose interactions determine the chain’s *fold* geometry, and a *side chain* that determines its amino-acid identity.

A protein’s *secondary structure* is its local pattern of α -helices and β -strands; its *tertiary structure* is the global three-dimensional fold those elements form when they pack together. A CATH code [63] (Class, Architecture, Topology, Homology) classifies this fold hierarchically. *Class* captures broad secondary-structure content (mostly- α , mostly- β , or mixed), *Architecture* the gross spatial arrangement of those elements, and *Topology* their specific connectivity.

For the structure-generation problem we study, it is customary to only model the backbone, with sequence identity recovered downstream using a separate *inverse-folding* model such as ProteinMPNN [24]. A further compression is the *CA-trace*, which uses only the set of C_α coordinates for every residue to represent the entire protein. This parameterisation retains enough information to recover fold topology and reduces computational cost relative to the more complete representation, which is typically referred to as *all-atom* modelling. Proteina [14] adopts this C_α parameterisation, representing each backbone as a sequence of per-residue tokens carrying continuous C_α Cartesian coordinates. The models reported in the original presentation are largely trained on high-confidence synthetic structures from the AlphaFold Database (AFDB) [36].

Once again, the CA-trace parameterisation is one of several techniques used in the literature to parameterise proteins in-silico. *Sequence-only* models, such as the ESM lineage [23], operate on amino-acid identity strings. These models capture rich functional priors, but encode far less information on 3D structure. *Full-atom* representations close the gap to physical realism, but add substantial compute cost. *Equivariant* parameterisations, such as torsion angles and oriented residue frames, encode SE(3) symmetries by construction, and form the standard basis for the equivariant generators noted above. However, these architectures have historically bottlenecked scaling. We refer the reader to Notin et al. [61] for a broader survey of machine-learning approaches to protein design.

2.5 Open questions

This survey leaves two questions open that the rest of this report takes on:

- (i) *Does REPA work in the molecular domain?* Existing applications of REPA in molecular generation [7, 8] have not been ablated over encoder choice, and protein backbone generation in particular remains unexplored.
- (ii) *When does REPA work?* The encoder-characterisation work surveyed in §2.2.1 has begun to ask which properties of an encoder predict REPA gain, but only within the image setting. We extend this question to the molecular domain.

We take both questions up empirically in Chapters 4–6.

Chapter 3

Evaluating molecular generation

What does it mean for a model to generate “good” artefacts? In this chapter, we outline a framework for evaluating molecular generation, and the tensions inherent to this endeavour. We argue that no single number suffices here, define the metric suites we use for small molecules (Table 3.1) and protein backbones (Table 3.2), and introduce the probe principle by which we measure a model’s internal representations separately from its outputs. We also introduce two themes that recur through this report in how we interpret these numbers. First, metrics can *saturate*, leaving no headroom for a method to register (Chapter 5). Second, metrics can form a *trade-off frontier*. A gain on one axis typically implies a loss on another (Chapter 6), so finding an acceptable balance, or better yet, pushing this envelope, is very useful.

3.1 What makes for a “good” molecule?

“Goodness” is usually task-specific, and a de novo design pipeline (§2.3) asks several things of a generated candidate pool at once. A useful set must at least be *valid*: realistic enough to survive downstream filters; *diverse*: spread across the design space, not collapsed onto a few modes; and *designable*: amenable to synthesis.

These axes present a real tension. A model can buy validity by memorising a handful of safe scaffolds, but at the cost of diversity. It can buy diversity by spreading its mass thinly across the learned distribution, but at the cost of designability. Maximising any one axis in isolation is uninformative, especially because these metrics are proxies for a sequence of filter steps in the canonical de novo design pipeline (§2.3). No axis is dispensable, yet these axes also trade off against each other.

3.2 Representation and generation are multi-factorial

Moreover, the need for a suite of evaluation metrics reflects how molecules are *represented* in-silico. This is in stark contrast with the image domain, which originated the generative machine learning methods we adopt (§2.1).

In a 2D image, *geometry* and *content* are coupled in one representation. The pixel grid is at once an object’s shape and its identity. There is a single contiguous feature space, so one self-supervised encoder can serve as a near-complete description of it. This coupling is what lets image generation lean on a single strong encoder, and lets image quality be summarised, to a first approximation, by a single Fréchet distance in that encoder’s feature space.

However, molecules pull geometry and content apart, and into more than two pieces. A molecule is at once a discrete graph of typed atoms drawn from a periodic-table-sized alphabet, and a continuous 3D geometry. This is a discrete–continuous duality with no clean analogue in the pixel grid. The two concepts are only loosely tied: one bonding graph admits many conformers, and similar geometries can host different atom types. Geometry is itself organised across scales, from local environments up to global shape. For example, with proteins the hierarchy is explicit, running from primary sequence through secondary-structure elements to the tertiary fold.

As a consequence, both representation and generation quality are multi-factorial in this domain. No single encoder captures content and geometry at every scale at once, as we see in Chapter 4. Moreover, no single metric can score a generated set along all those factors at once. A model can match the distribution of folds while getting local geometry subtly wrong; it can place atoms in valid positions while collapsing the diversity of shapes it produces. These are independent failures, so mode collapse is harder to reason about than in the image domain, where one number suffices to track what goes wrong.

The metric suite we propose in this chapter is an attempt at handling this complexity. We decompose quality along the factors the domain pulls apart, and interpret metrics holistically to understand the complete picture. For each domain we report two complementary families: *representation quality*, which scores what the model encodes internally, and *generation quality*, which scores its outputs. We present them in that order throughout the report, since representation is what REPA’s mechanism acts on (§2.2) and generation is the outcome we hope it improves.

3.3 The linear probe principle

One way to understand what a model represents internally is the *linear probe*. We freeze the model, extract its hidden states for clean inputs, and fit a simple linear regression to a known structural label. If the label is recoverable by this probe, that information is considered to be present and accessible in the representation; if not, it is absent or deeply entangled. This follows standard practice in image generation, where encoders are often characterised by ImageNet linear-probe accuracy [5]. We use this mechanism in two places: on the frozen encoders we might align to (Chapter 4), and on the generator trunk itself (Chapters 5 and 6).

The form of the read-out follows the label. For real-valued targets, we fit a regression and

score recoverability by held-out R^2 . For categorical targets, we fit a logistic classifier and score top-1 accuracy.

We outline the small-molecule targets we regress against in §3.4, and the protein suite in §3.5. Each probe maps to a necessary condition for a useful alignment target. The tasks span content and geometry from local to global scales, so we interpret a representation that scores well across the whole battery as genuinely carrying holistic information.

3.4 Small-molecule metrics

Table 3.1 collects every small-molecule metric we use, following the conventions of recent 3D molecular generators [15, 50, 52].

Representation quality. This asks what structural information the generator has learned to encode internally. We apply the probe principle of §3.3, regressing mean-pooled per-atom features onto four global RDKit descriptors [60] — molecular weight, LogP, ring count, and rotatable-bond count — and score recoverability by held-out R^2 . These descriptors span size, lipophilicity, topology, and flexibility, representing a wide range of chemical properties.

Generation quality. This asks if the generator produces good artefacts. For small molecules, quality metrics fall into three groups: per-molecule *structural quality* (validity, connectivity, uniqueness), set-level *coverage* (diversity, novelty), and *distributional realism* (the Fréchet ChemNet Distance, FCD [57]). The first group guards against malformed or repeated molecules; the second against mode collapse and memorisation; FCD, the small-molecule analogue of the image-domain Fréchet Inception Distance [58], asks whether the generated set as a whole resembles real molecules. As we explore in Chapter 5, the structural-quality metrics are particularly amenable to saturation by a strong generator. This does not mean the problem is solved; rather, it is a limit of the yardstick itself.

3.5 Protein-backbone metrics

Protein backbones are larger and more variable than small molecules (Chapter 2), and evaluating them is correspondingly more involved. Table 3.2 collects all the protein metrics we use. We present representation first, then generation, and close by offering our own framework for how the generation metrics should be grouped and read together.

Representation quality. We apply the probe principle of §3.3 for three tasks adapted from the literature, spanning local to global structure: per-residue inverse-folding recovery [27] for the local chemical environment; backbone dihedral error [26] for low-level geometry; and CATH fold classification [26] for global fold. The first two are scored per residue (top-1 accuracy and mean absolute error), while the fold probe is a logistic classifier on mean-pooled per-chain features. Because the three are largely orthogonal, improving all of them at once is strong evidence of holistic representation quality.

Metric	Definition	Content	Score
Representation quality — probe recoverability of RDKit descriptors [60]			
MolWt $\uparrow$	Molecular weight	Global size	R^2
LogP $\uparrow$	Octanol–water partition coefficient	Lipophilicity	R^2
Rings $\uparrow$	Ring count	Topology	R^2
RotBonds $\uparrow$	Rotatable-bond count	Flexibility	R^2
Generation quality			
<i>Structural quality</i>			
Validity $\uparrow$	Fraction sanitising to a valid molecule once bonds are inferred	Malformed geometry	Fraction
Connectivity $\uparrow$	Fraction forming a single connected component	Fragmentation	Fraction
Uniqueness $\uparrow$	Fraction not duplicating another generation	Trivial repetition	Fraction
<i>Set coverage</i>			
Diversity $\uparrow$	Mean pairwise dissimilarity across the generated set	Mode collapse	Fraction
Novelty $\uparrow$	Fraction of the set absent from the training data	Memorisation	Fraction
<i>Distributional realism</i>			
FCD $\downarrow$ [57]	Fréchet ChemNet Distance: Fréchet distance over the whole set in a property-prediction network’s feature space	Distribution mismatch	Distance

Table 3.1: Small-molecule metrics (Chapter 5), following the conventions of recent 3D molecular generators [15, 50, 52]. Arrows give the good direction.

Representation alignment. This asks not what the generator encodes, but how geometrically close it sits to a frozen encoder. We measure it with CKNNA [64], a mutual-nearest-neighbour kernel alignment, following the original REPA analysis [5] and a recent molecular extension [8]. We use it as a diagnostic lens on *how* REPA acts rather than a quantity we optimise.

Generation quality. It draws on two lineages of metric. The *canonical* ones — designability, and various counts of novelty and diversity — score or tally individual generations, and predate modern generative models. The *model-based* ones, introduced to this domain by Proteina [14], borrow the idea from image generation of comparing a generated set against a reference in a learned feature space. We headline on one from each: designability as the canonical measure of per-sample quality, and the Fréchet Protein Structure Distance (FPSD) as the model-based measure of set-level fidelity. We take each in turn, add a secondary-structure metric of our own, and then set out the framework that organises them all.

Designability. This metric estimates the fraction of generated backbones that are realisable as real proteins. It involves an in-silico round-trip through the protein design loop (§2.4.2). Each backbone is inverse-folded with ProteinMPNN [24], a model trained

Metric	Definition	Content	Score
Representation quality			
IF top-1 $\uparrow$	Per-residue inverse-folding sequence recovery [27]	Local chemical env.	Top-1 acc.
Dihedral MAE $\downarrow$	Per-residue backbone dihedral-angle error [26]	Low-level geometry	MAE
CATH-A top-1 $\uparrow$	Per-chain fold-architecture classification [26] (Table A.2)	Global fold	Top-1 acc.
Representation alignment			
CKNNA $\uparrow$ [64]	Per-residue mutual-nearest-neighbour kernel alignment, trunk vs. encoder	Trunk–encoder similarity	Score
Generation quality			
<i>Quality</i>			
FPSD $\downarrow$	Fréchet Protein Structure Distance: distance over the whole set in a structural encoder’s (GearNet) feature space	Distributional fidelity	Distance
Des $\uparrow$	Fraction of generated backbones that are designable	Physical realisability	Fraction
<i>Tertiary structure — whole set (T-W)</i>			
fJSD-A $\downarrow$	Jensen–Shannon divergence vs. reference CATH-A distribution	Fold-distribution fidelity	Divergence
fS-A $\uparrow$	Spread of generated folds across CATH-architecture classes	Whole-set fold diversity	Count
<i>Tertiary structure — designable subset (T-D)</i>			
#Clust $\uparrow$	Number of TM ≥ 0.5 structural clusters	Designable diversity	Count
pwTM $\downarrow$	Mean pairwise TM-score	Designable diversity	TM-score
<i>Secondary structure — whole set (S-W)</i>			
ssJSD-2D $\downarrow$	Jensen–Shannon divergence of the (helix-frac, strand-frac) distribution	Secondary-structure fidelity	Divergence
<i>Secondary structure — designable subset (S-D)</i>			
ssJSD-2D $\downarrow$	As above, on the designable subset	SS fidelity (designable)	Divergence
E%des $\uparrow$	Mean β -strand (E) fraction	β -content	Fraction

Table 3.2: Protein-backbone metrics (Chapter 6). Arrows give the good direction; full per-metric results appear in Tables 6.5 and 6.6.

to predict an amino-acid sequence that would fold into a given backbone, to propose candidate sequences. These sequences are then re-folded with ESMFold, a model trained to predict 3D structure from sequence. A backbone is termed designable when the self-consistency RMSD (scRMSD) between the generated backbone and its re-folded prediction falls below 2 Å. One caveat affects interpretation: designability is an in-silico proxy, not ground-truth physical realism, so it inherits the biases of the folding models it relies on. The inverse-folding and folding steps share model lineage with AlphaFold2 [45], and hence with the synthetic AlphaFold Database (AFDB), whose structures are themselves AF2 predictions. Generators trained on AFDB therefore score systematically higher on designability by construction, as we explore in Chapter 6.

FPSD. The Fréchet Protein Structure Distance [14] is the protein analogue of the imaging FID. First, it embeds the generated and reference sets using an external encoder: in this case, a GearNet [25] model trained on a fold classification task. Then, it models each set of embeddings as a Gaussian, and reports the Fréchet distance between the two — a scalar that drops to zero as the generated distribution matches the reference in both mean and covariance. Because the encoder reads fold-level geometry, FPSD scores how well a generated set reproduces the reference distribution of folds.

The complementary dynamic between designability and FPSD. Designability asks whether each backbone is individually realisable, while FPSD asks whether the set as a whole looks like real proteins. A model can do well on one and poorly on the other — collapsing onto a few highly designable folds, or spreading realistically while producing few foldable backbones — which is why we headline on both.

ssJSD-2D. We introduce this secondary-structure metric in our report. Like FPSD and fJSD, it is a *fidelity* metric: it scores how closely the generated secondary-structure distribution matches that of a real reference set. Proteina already measures secondary-structure fidelity, but only at the level of the mean helix and strand content, which is blind to a model that matches the average composition while collapsing its spread. In contrast, ssJSD-2D is a Jensen–Shannon divergence over the full joint (helix-fraction, strand-fraction) distribution against the reference, so it rewards reproducing the whole spread of secondary-structure content, not just its centre (Appendix A.6).

A two-axis taxonomy. We propose reading all these metrics through a framework that organises them by structural and distributional relevance. Two axes do the work: *tertiary* versus *secondary* structure, and the *whole set* versus the *designable subset*. The result is a grid of metric groups, which the result tables in Chapter 6 lay out as super-columns. Tertiary-whole (T-W) asks whether the distribution of global folds matches a reference over all backbones; its flagship is FPSD, alongside fJSD (a Jensen–Shannon divergence between the generated and reference fold-class distributions) and a fold-spread score (fS) [14]. Both fJSD and fS are defined at each level of the CATH hierarchy (§2.4.2); we headline on the architecture level (fJSD-A, fS-A). Tertiary-designable (T-D) asks the same diversity question, but only of backbones that pass designability, via cluster count

and pairwise TM-score. Secondary-whole (S-W) and secondary-designable (S-D) apply ssJSD-2D over the same two populations. (A Foldseek-based *novelty* axis sits alongside the five groups as an auxiliary measure. We report it where relevant, but never headline it.)

The tension between whole-set and designable-subset quality. The whole-set-versus-designable-subset split across structural hierarchies exposes a fundamental trade-off. A model can broaden the folds it generates *overall* (raising whole-set diversity, T-W) while concentrating its *designable* output onto a few reliable modes (lowering designable-subset diversity, T-D), and vice versa. Whole-set and designable-subset diversity are therefore distinct objects, and a method can move them in opposite directions. This is the designability-versus-diversity trade-off first noted in Proteina [14], and it is the thread we follow in Chapter 6.

3.6 Interpreting trade-offs to evaluate models

Ultimately, how do we use these metrics, which often trade-off against each other, to compare models? In our analysis, we distinguish between two phenomena. The first is whether a model merely occupies a different spot on the same trade-off envelope, in which case, which model is “better” is a use-case choice. The second is whether a model pushes the envelope itself. In this case, it plays the same trade-off game, but at a higher level, winning on one axis without conceding the other. This makes a model canonically better. We use this framework in the studies that follow, and in Chapter 6, we find evidence to support the latter, stronger claim in the context of REPA for protein generation.

We also note that generative metrics are noisy, so a gap between two point estimates is not necessarily a finding. We report bootstrap intervals and seed sweeps where possible to guard against this.

The complex dynamics of these metrics further motivate the question: *when can an encoder transfer useful representational information in this domain?* We examine this question in Chapter 4, which informs the experiments we present in Chapters 5 and 6.

Chapter 4

Encoder targets and profiling

What properties make an encoder a good target for representation alignment? In the image domain, the question remains unsettled, as we discussed in §2.2.1. Yu et al. [5] read *global* semantic quality as the driver; Singh et al. [6] instead find that the *local* spatial structure of an encoder’s token representations is what predicts REPA gain.

In this chapter, we ask this question in the molecular domain. Recall from §3.2 that molecular representation is *multi-factorial*: geometry, content, and structure across scales are pulled apart, and no single pre-trained encoder captures all of them. An encoder here is typically optimised for one task in isolation, and sees a molecule through that lens. An interatomic-potential model reads local geometry; a fold classifier reads global shape; a sequence model reads chemical identity. The encoder we align to therefore selects which factor REPA is most likely to transfer.

Our profiling work in this chapter makes this concrete. We build a suite of diagnostics to assess what information an encoder contains, how much of it is relevant to the generative task, and how well conditioned its embeddings are. We use this to profile seven candidate encoders across small molecules and proteins, and then use the results to guide which encoders to align with in the studies of Chapters 5 and 6.

4.1 A profiling protocol

We profile every candidate encoder with the protocol summarised in Table 4.1. Each encoder is run frozen on a sample of structures from our training sets: 200–2,000 molecules (QM9 and GEOM) for the small molecule encoders, and 200 proteins (PDB) for the protein encoders. This yields one embedding vector per atom or residue — tens of thousands of vectors in each case — and every diagnostic is computed from these per-token embeddings.

We ask three things of each encoder, and a good target needs all three.

Information asks: *what does the encoder contain?* This estimates the upper bound on what REPA can add to a model. Alignment cannot teach anything the encoder does

Diagnostic	What it captures	How we compute it
Information	<i>what the encoder encodes — an upper bound on REPA’s transfer</i>	
3D geometry	sensitivity to 3D pose	cosine self-similarity across conformers (molecules) or after a 1 Å coordinate perturbation (proteins)
Token identity	atom or residue type (high = a lookup the generator already has)	linear-probe accuracy
Relevance	<i>what the generator can match without improving its representation</i>	
Saturation	the embeddings barely vary, so a constant prediction already matches them	a high <i>floor</i> : the average cosine similarity of each token embedding to the dataset mean
Cheatability	the variation that exists is content the generator is already fed	a high <i>lift</i> from identity alone: the rise a REPA-form 3-layer MLP gains from a one-hot of the atom or residue type (cosine-similarity loss, 80/20 split)
Off-task risk	the remaining variation is some factor the generator does not need, rather than the 3D geometry it does	a low floor and low lift, read against the 3D-awareness probe
Conditioning	<i>whether the embedding is well-formed</i>	
Effective rank	dimensions actually used	RankMe [31] (Roy–Vetterli effective rank)
Sparsity, dead dims	inactive dimensions	fraction of exact zeros; permanently inactive dimensions
Norm scale	numerical conditioning	mean embedding ℓ_2 norm

Table 4.1: The profiling diagnostics, grouped by the three properties: information, relevance, and conditioning.

not itself hold. For a 3D generator, the axis that matters most is 3D geometry. We measure geometry-awareness by how far the embedding moves under a change in 3D pose: across distinct conformers of the same molecule for the molecular encoders, and after a 1 Å coordinate perturbation for the protein encoders, where only a single structure is available. We also report token identity, since a high score there suggests that an encoder is performing an identity lookup that the generator may already be capable of.

Relevance asks: *how much of a REPA target can the generator match without improving the representation it is trying to learn?* A generator does not need to learn all the information contained in an encoder. Indeed, during training, it is fed most of the relevant *content* of the molecule directly: the normalised token index always, and, for molecules, a one-hot of the atom type (the protein generator is given position but not residue identity; see §4.3). What it must acquire is the 3D *geometry*, which it sees only as noisy coordinates it is being trained to denoise. Hence, a REPA target is relevant only insofar as matching it forces the generator to improve that geometric representation.

A target can fail this in two ways:

- *cheatable*: the target is trivially reconstructible from inputs the generator already holds. So the model can saturate alignment loss without actually learning a better representation.
- *off-task*: matching the target does require representation learning, but what it requires is not what we need. For example, since the protein generator is not given residue identity, matching an identity-encoding target would force it to infer sequence from backbone. This is a genuine learning problem, but a different one from generating geometry, so the pressure may not assist the task we care about.

We estimate the cheatable fraction directly. We fit a three-layer MLP, the same projector that REPA uses (§2.2), to reconstruct embeddings from cheap features that withhold all 3D geometry: first *identity* alone (a one-hot of the atom or residue type), then *identity+position* (Table 4.1). Separately, with no projector fit, we also score the *constant* dataset mean, which is the optimal input-free prediction. We read this constant baseline as the *floor*, and the rise of the fitted rungs above it as the *lift*.

Read together, the floor and lift sort an encoder into one of three configurations:

- high floor suggests *saturation*. The embeddings barely vary, so a constant prediction already matches them. There is little signal to align to.
- low floor, high lift suggests *cheatable*. The embeddings vary, but the cheap inputs reconstruct that variation, so it is largely an identity lookup using data the generator is already fed.
- low floor, low lift is *promising*. However, a low lift only says the variation does not stem from the cheap identity we tested. It does not positively affirm that it comes from the geometry a generator needs. We interpret this against the information probes to tell a geometric encoder from one varying for some *off-task* reason.

Conditioning asks: *is the embedding mathematically well-formed enough to be a target?* For example, an embedding whose variance has collapsed onto a few directions (a low effective rank), or that is mostly zero, is a bottleneck. Too little structure passes through for the alignment to transmit anything, no matter the information or the relevance.

We apply this protocol to molecule encoders first (§4.2), then protein encoders (§4.3), and draw together what it implies in §4.4.

4.2 Small molecule encoders

We profile two publicly available pre-trained encoders, summarised in Table 4.2 (full values in Appendix A.8). They bracket a useful contrast. One sees the 2D bonding graph, the other sees 3D geometry. Neither, the framework finds, is a clean alignment target.

CheMeleon [30] is a directed message-passing network (a ChemProp D-MPNN) that passes

Encoder	Dim	token-ID	3D	Floor	Lift	Sparsity	Verdict
CheMeleon (2D graph)	2048	1.00	1.00	0.43	+0.04	93.8%	too sparse
MACE (3D potential)	192	1.00	0.998	0.86	+0.005	0%	saturated

Table 4.2: Molecule encoder profiles. Neither encoder shows promise. *token-ID*: atom-type linear-probe accuracy (trivially 1.00 for both; uninformative). *3D*: cosine self-similarity across conformers (lower = more 3D-sensitive). *Floor*: constant-prediction cosine (high = saturated). *Lift*: the rise a same-form MLP gains above the floor from cheap inputs. *Sparsity*: fraction of exact zeros.

Encoder	Dim	AA-probe	3D	Floor	Lift	RankMe	Verdict
GearNet	512	0.137	0.27	0.43	+0.009	256	usable
random-init	512	0.127	1.00	0.95	+0.003	100	saturated
ProteinMPNN	128	0.327	0.73	0.84	+0.014	85	weak but usable
ESM2-650M (L33)	1280	0.998	N/A	0.67	+0.053	1031	off-task [†]
PW-GearNet	3072	0.928	0.86	0.71	+0.013	208	poorly conditioned
MC-GearNet-Edge	3072	0.146	0.77	0.86	+0.002	12	collapsed

Table 4.3: Protein encoder profiles. GearNet and ProteinMPNN show promise. *AA-probe*: residue-identity linear-probe accuracy (high = an identity lookup the generator already has). *3D*: cosine self-similarity after a 1 Å perturbation (lower = more 3D-sensitive; N/A for sequence-only ESM2; moot for collapsed MC-GearNet). *Floor*: constant-prediction cosine (high = saturated). *Lift*: the rise a same-form MLP gains above the floor from cheap inputs. *RankMe*: Roy–Vetterli effective rank out of *Dim* [31]. [†]ESM2’s lift is the largest here but is sequence context, not 3D: off-task.

messages along the bonds of the 2D molecular graph; it never sees coordinates. Its embeddings are well spread (a low floor), but it fails on two other counts. First, it is 93.8% sparse, so alignment runs over a shifting, ill-conditioned subspace. Second, it is conformer-invariant by construction, so its lift is 2D bond-graph context, which may not aid a 3D generator.

MACE [29] is an equivariant message-passing network built on many-body atomic interactions, pretrained as an interatomic potential to predict energies from 3D structure. The model’s embeddings are mathematically well-formed. However, its energy pretraining captures local environments over global pose, so its 3D signal is weak, and a high floor of 0.86 suggests it is mean-saturated.

4.3 Protein backbone encoders

Proteins offer a wider choice of pre-trained encoders. We profile five candidates plus a random-initialised reference, summarised in Table 4.3 (full values in Appendix A.8).

GearNet [25], our lead candidate, is a message-passing network. It passes messages over a residue graph wired by spatial and sequential proximity. The checkpoint we profile is pre-trained for CATH fold classification and released with Proteina [14]. It is genuinely 3D-aware (a 1 Å perturbation drops the cosine self-similarity to 0.27) and barely a residue-

identity lookup (AA-probe 0.137). Its lift is small only because it sits on a low floor. The embeddings are spread, and the remaining headroom is likely geometric. The random-init row suggests that the architecture alone gives a near-constant, low-rank embedding, and the training significantly improves on this.

ESM2 [23] offers an instructive contrast, as a transformer-based protein language model trained by masked-residue prediction over sequence alone. It produces the largest lift in the field, yet that lift is sequence context. It ignores coordinates by construction, and its last layer flattens almost entirely to residue identity (AA-probe 0.998 against GearNet’s 0.137). This information is real but likely unhelpful to a backbone generator.

ProteinMPNN [24] is a message-passing network over a nearest-neighbour backbone graph, pretrained for inverse folding (predicting sequence from structure). It has the cleanest conditioning of any 3D-aware encoder we profile, but its geometric signal is significantly weaker than GearNet’s (a 1 Å perturbation leaves the cosine self-similarity at 0.73, against 0.27), and it has a relatively high floor. We characterise this encoder as weak but usable. Its 3D signal is faint, but it is clean and genuinely geometric.

The remaining two are GearNet-Edge variants. PW-GearNet [25, 26], pretrained on backbone torsions, is only weakly 3D-sensitive and poorly conditioned: its effective rank uses just 208 of 3072 dimensions. MC-GearNet-Edge [25] is collapsed entirely, with 95% of its variance in a single dimension.

Of the five candidates, only GearNet clears all three diagnostics, with ProteinMPNN also showing promise. Each of the others fails a different gate.

4.4 Informing our experiments

Our profiles are measured before any alignment training, so the expectations they set are genuinely a priori, not post hoc rationalisations. We use them to guide which encoders our studies will showcase. For small molecules (Chapter 5), we expect no measurable REPA gain. For protein backbones (Chapter 6), we expect that GearNet and ProteinMPNN are most likely to help. The following two chapters put these ideas to the test.

Chapter 5

REPA on small molecules

Previously (Chapter 4), we profiled various encoders in the molecular generation domain for their suitability as a representation alignment (REPA) target. This chapter is the first of two that put this analysis to the test. We begin here with the relatively simpler task of generating small molecules, before moving to much larger molecules in proteins next. In this chapter, we ask: *can REPA accelerate or improve the training of Tabasco [15], a flow-matching transformer-based generator of 3D small molecules?*

Our headline finding is that no REPA variant measurably improves, or accelerates the training of, Tabasco’s small-molecule generation in our setup. On structural-quality metrics, the baseline is already at ceiling. On FCD, the one metric with headroom, the REPA variants perform no better than baseline. Nor do they reach any level of quality faster: the baseline leads from the first epochs. We defer the *why* to §5.3, where two properties of this regime — a saturated evaluation surface and the absence of a usable teacher — emerge as the explanation and seed the contrast with proteins, which we explore in Chapter 6.

5.1 Experimental setup

5.1.1 Dataset

We train on GEOM-drugs [37], a dataset of drug-like molecular conformers at the sub-900-Dalton scale. It holds roughly 1.4M molecules, split $\sim 80/10/10$ into 1,142,099 training, 146,411 validation, and 146,798 test molecules. We model heavy atoms only, and augment each sample with eight random rotations. Each molecule is represented as a small point cloud of typed atoms, typically fewer than fifty, as outlined earlier (§2.4.1).

5.1.2 Models

The baseline is the unmodified Tabasco flow-matching transformer (§2.4.1): a 3.7M-parameter non-equivariant model that denoises 3D atom coordinates and atom types jointly. Our REPA variants add a trainable projector that aligns the trunk’s hidden

Property	Value
<i>Base architecture</i>	
Backbone (<i>trunk</i>)	Tabasco: 3.7M-param, 16-layer non-equivariant transformer (hidden dim 128, 8 heads, SiLU, cross-attention)
<i>REPA alignment</i>	
Encoder targets	CheMeleon (2D, 2048-d) · MACE (3D, 192-d)
Injection depth ℓ	final (last) layer
Combination mode	additive (alignment as plain regulariser)
Projector	single shared 3-layer MLP
Similarity metric sim	cosine similarity
Averaging mode	per atom
Alignment weight λ	0.8
<i>Training & sampling</i>	
Learning rate	2×10^{-3}
Batch size	256
EMA decay	0.999
Sampler	SDE (Langevin) interpolant, 100 integration steps
Prior noise scale	1.0 (unit Gaussian)
Langevin schedule	$1/(t+0.01)$, disabled for $t \geq 0.9$
White-noise scale	0.01

Table 5.1: Model and training configuration. This chapter presents two encoder targets.

states to a frozen encoder, evaluated on the clean ($t=1$) molecule (§2.2). The trunk carries two hidden-state streams, one for coordinates and one for atom types. We concatenate them and pass the pair through a single shared projector, so the alignment acts on one combined representation. It enters the loss as a plain additive regulariser, leaving the flow-matching objective untouched.

Encoder ablations. The two encoders we study come from the profiling of Chapter 4, and span a useful contrast. *CheMeleon* [30] is a 2D molecular-graph encoder. It is conformer-invariant, so it sees only the bonding graph and not the 3D pose. Meanwhile, *MACE* [29] is a 3D interatomic-potential encoder. It is geometry-aware, but, as the profiling showed, its output is a narrow descriptor of local atomic environments. Hence, the former carries rich 2D-graph information, while the latter encodes 3D geometry.

5.1.3 Metrics

Representation quality. We use the small-molecule probe battery of Chapter 3 (Table 3.1): a linear ridge regression from mean-pooled per-atom features to four RDKit descriptors [60], scored by held-out R^2 . A high R^2 means the descriptor is linearly decodable; an R^2 near zero (or slightly negative, when the fit predicts worse than the held-out mean) means that information is absent. We apply it to both the frozen encoders, to see what each REPA target carries on its own, and the generator trunk, to see what the model itself comes to encode, and rely on it in the diagnosis of §5.3.

Generation quality. We report the standard small-molecule generation metrics of Chapter 3 (Table 3.1): validity, connectivity, and uniqueness (structural quality); diversity and novelty (set coverage); and the Fréchet ChemNet Distance (FCD), the one distributional axis with headroom.

We compute every metric on a set of 1,000 molecules sampled from the trained model using the SDE interpolant described in Table 5.1. We mimic the molecular size distribution of GEOM-drugs in this set, and use the training set as the reference for our FCD metric.

The probe set comprises 1,000 molecules from the GEOM-drugs validation split. We fit the regression on 800 of these and report R^2 on the held-out 200.

5.2 Results: no measurable gain

No REPA variant meaningfully improves on the baseline. As outlined in Table 5.2, the baseline pushes validity, connectivity, and uniqueness to their ceiling, and the REPA variants follow suit. This is perhaps to be expected, since these are largely properties of bond topology, which a strong generator reproduces almost perfectly. Essentially, there is no headroom on these metrics for improvement.

On FCD, the one metric with room to move, both tested variants are worse. The baseline records the lowest (best) FCD at 5.61; CheMeleon lands at 7.43 and MACE at 6.81 (Table 5.2). We caveat this by noting that each variant is a single run, not a multi-seed average, and we do not have per-set FCD intervals to bound the gaps. Nevertheless, both point estimates favour the baseline.

There is no acceleration in generation quality. REPA’s headline result in the image domain is a training speed-up: the same generation quality reached in far fewer steps (§2.2) [5]. No such speed-up appears here. On every axis we can track over training, the REPA variants and the baseline are indistinguishable from the first few epochs (Figure 5.1). FCD, the one axis with headroom, is not logged over training, so we cannot trace its early trajectory; but its endpoint favours the baseline. REPA neither accelerates the saturated axes nor improves the one with room to move.

The unequal training does not make this comparison unfair. The baseline trained roughly twice as long as the REPA variants (~ 152 k steps over 33 epochs, versus ~ 69 – 73 k for the two REPA runs) for our final point estimates. We stopped the REPA runs early on purpose. With the quality metrics saturated and FCD flat or worse, there was no evidence that further training would change the verdict. Nor did the shorter runs cost the variants anything on the saturated axes, which plateau within the first few epochs for every run.




Figure 5.1: **REPA does not accelerate generation quality over the baseline.** Per-epoch generation quality on GEOM-drugs (100 molecules sampled per epoch during training), over the epochs all three runs reach. The baseline and both REPA variants track together and reach near-ceiling on every axis within the first epoch or two, leaving no headroom for alignment to fill. The metric is the in-training estimate, so values differ slightly from the final-checkpoint evaluation of Table 5.2.

Model	Steps	Valid.	Conn.	Uniq.	Div.	Nov.	FCD ↓
Baseline	151k	0.980	0.998	1.000	0.886	0.966	5.61
REPA-CheMeleon	73k	0.972	0.999	1.000	0.882	0.961	7.43
REPA-MACE	69k	1.000	0.995	0.996	0.884	0.977	6.81

Table 5.2: **REPA brings no measurable generation quality gain over the baseline.** The quality metrics are saturated, sitting at or near their ceiling for baseline and variants alike. On FCD, the only metric with headroom, both variants are worse. (1,000 sampled molecules; higher-is-better except FCD (↓). *Steps* is the training step at evaluation; the baseline trained $\sim 2\times$ longer, but the quality metrics plateau early, so the unequal steps do not make the comparison unfair (§5.2). One run per variant.)

5.3 Discussion

The null result suggests this regime offers neither headroom nor usable teachers. Under the original hypothesis, REPA [5] alleviates a representation bottleneck and therefore accelerates improvements on generation quality. However, this needs two conditions. First, there must be *headroom*: a gap in what the generator represents or produces, for the alignment to fill. Second, there must be a *usable teacher*: an encoder that encodes representational information that differs from what the trunk already has and transfers to the generative task. In this regime, neither condition holds.

The first condition fails on the metric surface. The generation quality axes are saturated, as §5.2 showed.

The second condition fails on both encoders, for opposite reasons (Table 5.3). We use linear probes to determine the representational content each frozen encoder carries against what the trunk already has. CheMeleon decodes the descriptors almost perfectly in isolation (R^2 up to 0.99). But the baseline trunk *already decodes them well* (0.64–0.92), because they are 2D-graph quantities and the trunk sees the graph. CheMeleon’s signal

Source	MolWt	LogP	#Rings	#RotBonds
<i>Frozen encoders (the REPA targets in isolation)</i>				
CheMeleon (2D)	0.993	0.988	0.995	0.989
MACE (3D)	0.022	0.307	0.081	-0.014
Dummy	0.822	0.160	0.432	0.345
<i>Generator trunk (the model’s internal representation)</i>				
Baseline	0.920	0.713	0.742	0.636
REPA-CheMeleon	0.943	0.669	0.736	0.606
REPA-MACE	0.935	0.703	0.704	0.582
REPA-MACE (tradeoff)	0.955	0.730	0.732	0.666

Table 5.3: **Linear-probe R^2 reveals little representational difference between REPA and the baseline.** The baseline trunk already captures what CheMeleon learns, while MACE’s embedding is too narrow to decode these descriptors. Hence, neither makes for a usable REPA target. (R^2 for four RDKit descriptors, on held-out molecules; §5.1.3.)

therefore overlaps with what the model already has. Moreover, FCD ultimately scores the inferred 2D graph — the very information the trunk already holds — so even on the one metric with headroom, CheMeleon adds nothing new.

MACE presents the contrary problem. It is genuinely 3D, but it encodes only a narrow, energy-tuned slice of *local* atomic chemistry. The global descriptors we probe sit outside that slice, so they are largely undecodable from its mean-pooled embedding (R^2 near zero, save a weak 0.31 on LogP). Its signal differs from the trunk’s, but does not transfer. To calibrate the scale, note that an untrained random-weight encoder (Dummy) already decodes molecular weight at $R^2 = 0.82$, simply because weight tracks atom count. MACE falls below even this floor on three of the four descriptors.

Forcing the signal in does not rescue it either. As a final test, we replaced the additive regulariser with a convex *tradeoff* loss, $(1 - \lambda) \cdot \mathcal{L}_{\text{gen}} + \lambda \cdot \mathcal{L}_{\text{REPA}}$. This deliberately down-weights the flow-matching loss to force the model to adopt the encoder’s representation. Even so, the effect on generation is negligible. MACE under the tradeoff loss still trails the baseline on FCD (6.32 versus 5.61). Its only mark is on the representation, where it is the only variant to beat the baseline on a majority of the descriptor probes (Table 5.3), but only by a small margin. Even when forced into the architecture, MACE leaves only a tiny trace in representation, and nothing in generation.

We do not claim to have exhausted the REPA design space here. We vary the encoder and the loss form, but hold the alignment layer and projector architecture fixed at the image-domain defaults (§2.2). The diagnosis below should therefore be read as the most parsimonious account of the regime, consistent with the encoder profiles of Chapter 4, rather than a proof that no configuration could help.

When is REPA actually useful? Our results sharpen the question on REPA’s effectiveness rather than settling it. Under our setup, REPA does not help small-molecule generation, because the regime offers neither headroom nor a usable teacher. This is not

evidence that REPA cannot help with molecular generation in principle. But it does suggest the question has no single answer. Molecular generation is a broad domain of distinct regimes, and whether REPA helps is decided within each, not across the whole. Read that way, its utility points less to a single intervention than to a family of them.

More practically, we must look for a regime within the domain that satisfies these two conditions. Protein backbone generation presents an intriguing possibility on both fronts, and that is where we turn next.

Chapter 6

REPA on protein backbones

In this chapter, we expand the scope of our study on representation alignment from small molecules to much larger ones: proteins. Previously (Chapter 5), our efforts using REPA to generate small molecules yielded little demonstrable improvement, because the Tabasco [15] baseline already saturated most measures of generation quality, and because the encoders used were likely unhelpful. However, de novo backbone design represents a more challenging task, as the data distribution of proteins is significantly larger, and the geometry of these structures is correspondingly more complex. In this context, we ask: *can REPA accelerate or improve the training of Proteina [14], a flow-matching transformer-based protein backbone generator?*

Our headline finding is that REPA accelerates convergence on the two key measures of protein generation quality. This holds across both experimental and synthetic training-data regimes (Figure 6.1). The one exception is designability on synthetic data, which the baseline already saturates under standard sampling (§6.1.1). This chapter characterises *what* happens and attempts to understand *why*. We also pick up the thread from Chapter 3 on how the molecular domain differs from the image domain REPA was built for, and develop a theme throughout: REPA here is a *family* of interventions rather than a single mechanism. Throughout, our aim is to test whether the mechanism transfers at all. We do not search for its best configuration (§6.2.6).

6.1 Experimental setup

We integrate REPA into Proteina [14], a 60M-parameter non-equivariant flow-matching generator of C α protein backbones (§2.4.2). Following §2.2, we attach a trainable projector at a chosen layer of the generator trunk (the transformer stack) that aligns the model’s hidden states to a frozen external encoder.

6.1.1 Datasets

We train in two data regimes that bracket the data-scarcity problem motivating this report (Chapter 1). The Protein Data Bank (PDB) [35] is the experimental record:




Figure 6.1: **REPA accelerates FID in both data regimes, and designability on experimental (PDB) data.** On synthetic (AFDB) data the designability baseline is already saturated, so we make no acceleration claim there (§6.1.1). The random-encoder control (grey) helps distinguish gains induced by regularisation alone from those induced mechanistically by an encoder’s learned structure. While PDB retains a persistent designability gap, its baseline appears to catch up on FID after $\approx 1\text{M}$ steps; we examine this in §6.2.3. ($n \leq 256$ models, 3 seeds; shaded bands show min/max.)

high-quality and ground-truth, but *scarce*, on the order of hundreds of thousands of structures. The AlphaFold Database (AFDB) [36] is the synthetic alternative: hundreds of millions of AF2-predicted structures, *abundant* but model-output and helix-biased. The field’s response to scarcity is precisely to substitute synthetic scale for experimental quality, so studying REPA in both regimes instruments that substitution directly; we treat them as two ends of a quality–scale axis, not a gold standard and a compromise. (We use the manually-curated Swiss-Prot subset of AFDB as a reproducible stand-in for the Foldseek-clustered TrEMBL subset on which Proteina was trained in the original presentation [14]. This corpus contained data that has since been pruned from AFDB, and is hence inaccessible.)

Following Proteina [14], we train on the subset of chains with at most 256 residues; we represent this length-regime by $n \leq 256$. (We note that this represents only a property of the training datasets, not a bound on the length of the proteins that models trained on these datasets can generate.) Following this filter, the two training sets are comparable in size (PDB: 273,771 / 3,190 / 323 train/val/test chains; AFDB: 229,670 / 4,521 / 235).

Appendix A (Table A.1) contains further details on dataset composition. During training, we augment each backbone with a single random global rotation per sample.

Different datasets in this domain carry different biases. Trends in designability offer a clear example, as training on AFDB raises the in-silico-designability floor well above PDB at every setting we tested (§6.2.3). This is the designability-proxy lineage effect of Chapter 3: the test re-folds with models sharing AlphaFold2’s lineage, so AlphaFold-like backbones are plausibly re-folded most confidently. A second caveat concerns leakage: the PDB random split has substantial train-validation sequence overlap, which inflates in-distribution scores unless a clean split is used (Appendix A.3). We flag both so the metrics in this chapter are read with care.

6.1.2 Models

The base Proteina model comprises a stack of ten transformer layers (zero-indexed) that feeds a flow matching objective. Our experiments add a REPA loss to this design, as outlined in Table 6.1.

Encoder and injection-depth ablations. Our narrative concentrates on the two encoders identified as most viable in Chapter 4: *GearNet* [25], a structure encoder pre-trained for CATH fold classification (checkpoint released with Proteina [14]); and *ProteinMPNN* [24], which encodes each residue’s local inverse-folding environment. Each encoder is attached to the trunk in two layer-depth configurations: layer 4 (middle layer) and layer 9 (last layer). The remaining encoder and depth combinations are collected in Appendix A. Figures presented in this chapter will typically feature the best performing model to serve as an illustrative example.

Separating regularisation from mechanism. A random-weights GearNet attached at layer 4, identical in architecture but never trained, serves as a control architecture in this chapter. We use this model to separate whatever REPA achieves merely due to the regularisation effect of adding an auxiliary loss from what it achieves through the encoder’s *learned* structural knowledge. Crucially, a trained encoder at the same layer 4 still outperforms the random one, so the learned-versus-random gap is not an artefact of injection depth (Appendix A).

Hyperparameters. Our defaults largely follow the original image-domain REPA recipe [5].

6.1.3 Metrics

The original hypothesis that motivated REPA for imaging [5] was that alleviating a *representation* bottleneck would improve *generation* quality. We seek to understand if the same mechanism applies in the molecular domain as well. Hence, along with estimating the quality of a model’s generated samples, we are also interested in what the model comes to *represent* internally. To this end, we measure three families of metrics, as outlined below: representation quality, representation alignment, and generation quality (Table 3.2).

Property	Value
<i>Base architecture</i>	
Backbone (<i>trunk</i>)	Proteina (CAFlow): 60M-param, 10-layer non-equivariant C α transformer (token dim 512, 8 heads)
<i>REPA alignment</i>	
Encoder targets	GearNet (512-d) · ProteinMPNN (128-d) · random-GearNet (control)
Injection depth ℓ	layer 4 (middle) · layer 9 (last)
Projector	3-layer MLP, hidden dim 512
Similarity metric sim	cosine similarity
Averaging mode	per residue
Alignment weight λ	0.5
<i>Training & sampling</i>	
Learning rate	10^{-3}
Batch size	24 ($n \leq 256$)
Weight decay	none
Sampler	SDE, noise scale $\gamma = 0.45$
Integration steps	400 (step size $dt = 0.0025$)
Self-conditioning	on
Diffusion schedule	$g_t = 1/t$, log-spaced time steps

Table 6.1: Model and training configuration. This chapter presents two encoder targets and two injection depths. Other ablations are deferred to Appendix A.

Representation quality. We use the three-task protein battery of Chapter 3 (inverse folding, dihedral error, CATH fold classification; Table 3.2). Because the three probe largely orthogonal capabilities — local chemical environment, local geometry, and global fold — improving all of them at once is strong evidence of holistic representation quality. Among the CATH levels we headline on architecture (CATH-A): Class is too easy (top-1 saturates) and Topology too sparse (too many buckets, too few chains). We report all three for completeness, the patterns holding across them, and motivate the choice in Appendix A (Table A.2).

Representation alignment. This measures how close the trunk sits to a structural encoder, via CKNNA (§3.3; protocol below). As in the original REPA analysis [5] and a recent molecular extension [8], we use it as a diagnostic — a mechanistic lens on *how* REPA might produce its representation gains — rather than a quantity we optimise or build claims on.

Generation quality. This is what ultimately matters, and we report it through the two-axis taxonomy of Chapter 3 — tertiary/secondary structure crossed with whole set/designable subset — laid out as supercolumns in our result tables: per-sample quality, then T-W, T-D, S-W, and S-D (Table 3.2). Two headlines anchor the rest: designability and FPSD. The whole-set-versus-designable-subset split, and the fidelity-versus-diversity reading that cuts across it, recur throughout this chapter.

Most of these metrics follow Proteina [14]; the one we add, ssJSD-2D, is defined in Chap-

ter 3 (full description in Appendix A.6). As we explore in §6.2.4, REPA improves fidelity and whole-set diversity together, but may trade off tertiary-structure diversity *within* the designable subset.

6.1.4 Evaluation protocol

Representation probing. Following the probe principle of §3.3, we fit probes to a frozen checkpoint’s hidden states on clean inputs ($t=1$): the inverse-folding and dihedral probes per residue, the CATH probe per chain on mean-pooled representations. IF is a 20-way classification, CATH-A a ~ 40 -way one.

Each probe is fit on 1,000 training chains. We evaluate the per-residue probes (inverse folding and dihedral) on a *blinded* set of ~ 325 proteins ($\sim 43\text{k}$ residues) held to under 30% sequence identity to all training data, which removes both model- and probe-side leakage. The per-chain CATH probe needs many distinct chains spread across fold classes, which that small slice cannot supply, so we relax it to a *probe-clean* evaluation on a larger validation set of $\sim 3,190$ chains. We detail this scheme, and the model/probe leakage distinction it controls for, in Appendix A.3, and read rank-order and Δ -from-baseline throughout this chapter.

Representation alignment. Recall that CKNNA scores how far two representations share their nearest-neighbour structure (§3.3). We compute it on clean ($t=1$) data from a fixed held-out batch, between the hidden states at every trunk layer and three domain-relevant encoders: GearNet, ProteinMPNN, and ESM2 [23]. We use $k=10$ neighbours over 10,000 residues, and bracket each value with a subsample-without-replacement bootstrap (50 draws at 80%; sampling with replacement would inflate the score through duplicate neighbour pairs). The absolute CKNNA values we record are small, an order of magnitude below image-domain REPA [5], so we read the direction and layer structure of the effect, and its separation from the baseline, not its raw magnitude.

Generation sampling. Our generation evaluation follows the Proteina protocol [14], at a reduced sample budget.

For whole-set evaluations, we sample 1,125 backbones per checkpoint — 125 at each length in $\{50, 75, \dots, 250\}$ — and compute metrics over all of them. Depending on the training regime, we use matching PDB or AFDB reference statistics.

For designability evaluations, 50 backbones at each length $\{50, 100, 150, 200, 250\}$ are inverse-folded with ProteinMPNN (8 sequences each), re-folded with ESMFold, and called designable when the self-consistency RMSD is below $2 \text{ \AA}$. The designable-subset metrics (T-D, S-D) are then computed on the backbones that pass.

With this experimental testbed in place, the rest of the chapter investigates what the REPA construction changes with respect to baseline, and how that helps improve generation quality. We now turn to the results.




Figure 6.2: **REPA offers a mechanistic effect that lifts the trunk’s representation quality asymptotically.** REPA variants open a quality gap over baseline that never closes. The random control gains least, evidence that a trained encoder contributes a genuine mechanistic effect rather than mere regularisation from the auxiliary loss. ($n \leq 256$, PDB-trained; best-layer linear probes; per-residue probes on a blinded cross-database slice, CATH on a probe-clean split; 3 seeds, shaded bands show min/max. See Appendix A.3 for further details and ablations.)

6.2 Results

6.2.1 REPA improves representation quality asymptotically

At baseline, the flow-matching objective helps the model encode some useful representational information. We probe our models along the three orthogonal axes of §6.1.3 — local chemical environment (inverse folding), local geometry (backbone dihedrals), and global fold (CATH classification, with CATH-A as our headline metric; Figure 6.2). The trained baseline consistently does better than both random features and an untrained ProteinA trunk. This suggests that the flow-matching objective does indeed learn to *represent* protein backbones as it learns to *generate* them.

Every trained-encoder variant improves representational capacity, consistently and durably. Each lifts all three probes substantially above the baseline, from the very beginning of training. That difference then persists. This is an asymptotic gap that the baseline never closes across 1.8M steps, not just a head start or an accelerated convergence pattern. This gap confirms one of the hypotheses posited in the original REPA presentation [5]: an alignment target during training can fill the representational headroom left by the flow-matching objective.

Every encoder routes representation gain differently. REPA-GearNet consistently wins the fold probes, while REPA-MPNN wins the per-residue probes (Table 6.2). This pattern follows what we would expect a priori, given the tasks that each encoder has been specifically trained for. This ties into our emerging theme on the multi-factorial nature of representation learning in the molecular domain (Chapter 3). No single encoder captures every factor, so the choice of alignment target is also a choice over the representational

Variant	IF top-1 $\uparrow$	dihedral MAE ($^\circ$) $\downarrow$	CATH top-1 acc. $\uparrow$		
			C	A	T
Baseline	0.130	27.6	0.717	0.397	0.303
REPA L4-random	-0.002	+2.7	+0.051	+0.029	+0.017
REPA L9-GN	+0.022	-3.6	+0.163	+0.283	+0.462
REPA L9-MPNN	+0.027	-6.8	+0.085	+0.078	+0.113

Table 6.2: **REPA offers persistent and encoder-routed representation advantages.** Consistent with what each encoder was trained for, REPA-GearNet wins the fold probes (CATH C/A/T), while REPA-MPNN wins the per-residue probes (IF, dihedral). We headline on CATH-A, and report C and T for completeness. ($n \leq 256$ PDB-trained; best-layer linear-probe performance, averaged over the 700K–1.2M training-step window; Δ from baseline; green / darker = improvement/best per probe.)

factor a model designer wishes to emphasise. We establish this routing on PDB-trained models and report the AFDB results in Appendix A.

REPA offers a mechanistic effect beyond just regularisation. Our random-weights control improves on the baseline for the fold probes (CATH), but not for the per-residue ones. Where it helps, this may be because *any* regularisation generically improves representation capacity. It may also be that even an untrained message-passing network imposes some geometric priors that aid the trunk. However, the trained encoders consistently perform far better. This sharpens the case for a genuine mechanistic effect, as regularisation alone is not the whole story.

The random control behaves exactly as the framework predicts a saturable target should. In §4.3 we characterised the random-init encoder as a near-constant, low-rank embedding, the high-floor signature our relevance criterion flags as *cheatable*: a target the trunk can match without learning anything (§4.1). The training loss bears this out. The random encoder’s alignment cosine similarity *saturates higher* than the trained one’s — ≈ 0.89 against ≈ 0.80 for a trained GearNet at the same layer 4 on PDB — precisely because its barely-varying target is trivially matched. Yet this near-perfect alignment buys the least representation quality of any variant, and no FPSD or designability gain (§6.2.3). A target can be easy to align to and still teach the model nothing, which is the failure mode our Chapter 4 framework was built to anticipate.

6.2.2 A diagnostic check on alignment

Following Yu et al. [5] and recent in-domain work [8], we investigate the mechanism that may be driving REPA’s representation quality gains by quantifying the representation alignment it induces. Using CKNN [64], we score how closely our models’ per-residue geometry matches that of three domain-specific encoders: GearNet and ProteinMPNN, our two alignment targets in this chapter; and ESM2 [23], which we align to only in our ablations, not in the models presented here. We note that we study alignment only as a diagnostic, and not to make claims about REPA’s utility.




Figure 6.3: **REPA aligns the trunk with domain encoders, including those it was not trained with.** (a) Each REPA variant lifts per-residue alignment to its own target encoder above the baseline floor. (b) REPA induces cross-encoder alignment: for example, a GearNet-aligned trunk also moves toward ProteinMPNN and ESM2. Absolute alignment is small throughout, so we read the pattern, not its magnitude. ($n \leq 256$ PDB-trained, step 1M, $t=1.0$; per-residue CKNNA, $k=10$; lines are bootstrap medians, shaded bands are the 5–95% subsample bootstrap interval.)

REPA lifts alignment above the random floor, albeit at low magnitude. The baseline sits at the random floor at every layer (CKNNA ≈ 0), while REPA separates from it cleanly, suggesting that its hidden states do move towards the encoder’s (Figure 6.3a). This serves as a minimal check that the construction changes model behaviour. However, the magnitude of this signal is small. Peak per-residue CKNNA (bootstrap median) is ≈ 0.02 – 0.05 , on a scale where self-alignment is 1. We therefore read the pattern, not its size, and stake no strong claim on these numbers.

Aligning to one encoder nudges the trunk towards others as well. Every REPA variant we tested raises alignment to all three encoders, including those it was not aligned to (full model $\times$ encoder matrix in Table A.3, Appendix A.4). A GearNet-aligned trunk, for example, also moves closer to ProteinMPNN and ESM2 [23] (Figure 6.3b), suggesting that REPA-induced alignment is not merely circular self-alignment.

One possible explanation for this phenomenon is the *Platonic representation hypothesis* [64]. This idea suggests models trained on related data converge toward a shared representation. In this framework, REPA may be nudging the trunk toward a common structural representation that several encoders approximate, rather than toward any single encoder.

6.2.3 REPA accelerates generation quality by climbing the representation–generation envelope faster

In this subsection, we focus on the holistic *Quality* metrics of our suite, FPSD and designability (outlined in Table 3.2), to outline our headline results. A deeper analysis of generation quality along tertiary- and secondary-structure decomposition follows in §6.2.4.

Variant (last ckpt)	Accel. @400K (vs base)	Long run (abs. best)
<i>AFDB FPSD</i> ↓		
baseline (to 1.7M)	431 (@400K)	386 (@1.3M)
GearNet-L4 (1.2M)	+24% (328)	282 (@1.2M)
GearNet-L9 (900K)	+28% (313)	313 (@400K)
MPNN-L4 (1.1M)	−10% (477)	416 (@1.0M)
MPNN-L9 (1.5M)	+18% (352)	352 (@400K)
Random control (500K)	+18% (353)	353 (@400K)
<i>AFDB designability</i> ↑		
baseline (to 1.7M)	0.75 (@400K)	0.75 (@400K)
GearNet-L4 (1.2M)	+4% (0.78)	0.78 (@400K)
GearNet-L9 (900K)	−4% (0.72)	0.77 (@800K)
MPNN-L4 (1.1M)	−9% (0.68)	0.69 (@700K)
MPNN-L9 (1.5M)	−6% (0.71)	0.77 (@700K)
Random control (500K)	+9% (0.82)	0.82 (@400K)
<i>PDB FPSD</i> ↓		
baseline (to 2.4M)	472 (@400K)	288 (@1.6M)
GearNet-L4 (1.0M)	+7% (440)	432 (@1.0M)
GearNet-L9 (1.6M)	+28% (341)	319 (@700K)
MPNN-L4 (1.6M)	−8% (512)	329 (@1.6M)
MPNN-L9 (2.0M)	+13% (410)	297 (@1.6M)
Random control (1.2M)	−7% (506)	351 (@1.1M)
<i>PDB designability</i> ↑		
baseline (to 2.4M)	0.26 (@400K)	0.67 (@1.6M)
GearNet-L4 (1.0M)	+98% (0.52)	0.70 (@700K)
GearNet-L9 (1.6M)	+97% (0.52)	0.63 (@800K)
MPNN-L4 (1.6M)	+124% (0.59)	0.71 (@1.6M)
MPNN-L9 (2.0M)	+149% (0.66)	0.82 (@1.8M)
Random control (1.2M)	+48% (0.39)	0.46 (@900K)

Table 6.3: **REPA accelerates generation quality over the baseline.** In some regimes it also wins the long run. Acceleration is the %-delta vs the baseline at 400K (the anchor of Figure 6.4), with the absolute score in parentheses; the long run is each variant’s absolute best within a common 2.0M-step window, coloured against the baseline-best row. Generation scores are means over up to three sampling seeds; $n \leq 256$; green / red = better/worse than baseline.

REPA gives generation quality an early head start. At 400K steps, every variant already leads the baseline on PDB designability, by +48% to +149%. Most variants also lead on FPSD in both regimes, with the best at +28%. AFDB designability is the lone exception, for the proxy reason given in §6.1.1 (Table 6.3).

This acceleration is consistent with the representation-bottleneck hypothesis. Yu et al. [5] hypothesise that a model’s representation quality and its generation quality are coupled, because it is learning both at once. In our PDB-trained models, this coupling is clear. Across checkpoints, representation and generation quality are positively correlated, and the strongest of these correlations survive controlling for training step (Table 6.4). The same coupling holds on AFDB-trained models, and in the same




Figure 6.4: **At equal compute, REPA achieves better representation and generation quality.** The dashed arrow marks the matched-compute comparison at step 400K. Against the baseline, the REPA variant has higher fold accuracy (CATH-A) and better generation lower FID (a) and higher designability (b). (Markers are a coarse $\approx 400\text{K}$ -spaced subset of checkpoints from 100K, with the 400K point bordered; x = student CATH-A accuracy, higher to the right; generation values are means of $N=3$ sampling seeds; $n \leq 256$, PDB-trained.)

encoder-matched form (Appendix A.5). If the two move together, then securing good representations early should secure good generation early as well.

REPA climbs the representation–generation envelope faster than the baseline. REPA reaches an asymptotically higher representation quality almost as soon as training begins (§6.2.1, Figure 6.2). On the representation–generation curve posited by the bottleneck hypothesis, this places it further along than the baseline at matched compute (Figure 6.4, shown for the strongest PDB variant). Hence, it reaches a given level of generation quality at less training. We interpret REPA as simply traversing this curve faster, rather than finding a new shortcut.

This representation–generation link is encoder-matched, suggesting the effect is encoder-specific, not generic regularisation. Each encoder’s generation gain lands on the axis where its representations are strongest. GearNet, the better fold encoder (§6.2.1), gives the larger gain in fold *distribution* (Table 6.5). MPNN, the better local encoder, gives the larger acceleration in *designability* (Table 6.3). Such axis-matching would be unlikely under a generic auxiliary loss, but we do not claim a formal proof of causality.

Broadly, REPA’s lasting advantage tracks the headroom the baseline leaves. Where the baseline plateaus poorly, REPA’s lead is large and permanent, as on AFDB FPSD. Where the baseline is a slow learner, the lead narrows yet still holds, as on PDB designability. Only where the baseline is a genuinely strong learner, as on PDB FPSD, does it eventually close the gap. AFDB designability, again, leaves little headroom over the baseline. However, as we noted earlier (§6.1.1), this is likely a quirk of the relationship between how designability is measured and how AFDB was generated, and not due to

Representation quality	FPSD		Designability	
	partial	raw	partial	raw
Inverse folding (top-1)	+0.37	+0.17	+0.79	+0.61
Dihedral MAE	+0.44	-0.03	+0.77	+0.36
CATH-A	+0.22	+0.11	+0.49	+0.38

Table 6.4: **Better trunk representations track better generation — strongly for designability, moderately for FPSD.** Partial Pearson correlation controlling for training step, with the raw value alongside; oriented so a positive value means better representation accompanies better generation. ($n=59$ checkpoints pooled over the baseline, the REPA variants, and the random control; PDB.)

Variant (700K)	Quality		T-W		T-D		S-W	S-D	
	Des $\uparrow$	FPSD $\downarrow$	fJSD-A $\downarrow$	fS-A $\uparrow$	#Clust $\uparrow$	pwTM $\downarrow$	ssJSD2D $\downarrow$	ssJSD2D $\downarrow$	E%des $\uparrow$
<i>PDB-trained, $n \leq 256$, 700K, 3 seeds; FPSD is FPSD-PDB (in-distribution).</i>									
Baseline	0.46	437	1.11	5.48	77	0.27	0.35	0.40	0.22
REPA L4-random	-0.08	+34	-0.55	+0.05	-23	-0.10	-0.13	-0.01	-0.11
REPA L4-GN	+0.24	+71	-0.38	+1.63	+1	+0.11	-0.21	-0.12	+0.00
REPA L9-GN	+0.14	-118	-0.57	+2.31	+27	-0.02	-0.19	-0.10	-0.04
REPA L4-MPNN	+0.18	-88	-0.13	+0.48	-4	+0.08	-0.16	-0.14	+0.01
REPA L9-MPNN	+0.05	-19	-0.32	+0.37	+33	-0.01	-0.14	-0.11	-0.07
<i>AFDB-trained, $n \leq 256$, 700K, 2-3 seeds (MPNN-L4: 1 seed); FPSD is FPSD-AFDB.</i>									
Baseline	0.72	494	2.45	4.59	128	0.22	0.15	0.16	0.15
REPA L4-GN	+0.05	-195	-1.56	+2.27	-35	+0.04	-0.11	-0.07	-0.02
REPA L9-GN	+0.01	-172	-1.29	+2.76	-50	+0.11	-0.12	-0.06	+0.04
REPA L4-MPNN	-0.03	-22	-0.39	+0.03	-1	-0.07	-0.02	+0.05	-0.01
REPA L9-MPNN	+0.05	-139	-0.93	+0.31	+21	-0.03	+0.01	+0.02	-0.04

Table 6.5: **Across the metric suite, almost every REPA variant improves whole-set fidelity and diversity over baseline.** The effect is broad enough that even the random-weights control matches the reference distribution better. However, only the learned encoders also lift per-sample quality (FPSD and designability). ($n \leq 256$, step 700K; baseline absolute, REPA rows show Δ ; green / red = improvement/regression; header arrows mark the good direction. Whole-set diversity is read from fS-A; designable diversity from #Clust and pwTM.)

any special merit. Tellingly, even the random control scores highly here.

6.2.4 REPA improves whole-set coverage but may trade off diversity within its designable subset

In this subsection, we investigate how REPA impacts different aspects of generation quality beyond our headline metrics. We read these observations as interesting qualifiers that provide further context to the main convergence result discussed above (§6.2.3).

REPA consistently improves whole-set fidelity and diversity metrics. Under almost every REPA variant tested, the generated set matches the reference fold distribution more closely (fJSD-A) and spreads more widely across fold classes (fS-A); secondary-structure fidelity (ssJSD-2D) improves likewise. Curiously, this is true of the random-weights control as well. However, only the learned encoders translate this advantage into

Designable pwTM ↓	700K steps			1.0M steps		
Variant	$\beta < 10$	10–25	$\beta \geq 25$	$\beta < 10$	10–25	$\beta \geq 25$
<i>PDB-trained</i>						
Baseline	0.15	0.22	0.50	0.15	0.23	0.13
REPA L9-GN	0.16	0.27	0.73	0.14	0.24	0.87
REPA L9-MPNN	0.16	0.21	0.67	0.26	0.22	0.67
<i>AFDB-trained</i>						
Baseline	0.14	0.29	0.17	0.16	0.26	0.15
REPA L9-GN	0.37	0.31	0.82	0.40*	0.32*	0.76*
REPA L9-MPNN (<i>falsifier</i>)	0.15	0.27	0.11	0.15	0.24	0.15

Table 6.6: **A REPA-learned distribution may collapse onto fewer distinct folds within its designable subset.** The effect is strongest in the β -strand-rich folds. (Designable-subset pairwise TM, lower = more diverse; columns bin the designable subset by β -strand fraction (%); colour marks REPA *vs.* baseline at the same step and bin: red / green = more concentrated / more diverse, $|\Delta| < 0.05$ uncoloured. $n \leq 256$, single-seed; *900K steps.)

higher FPSD and designability (Table 6.5).

However, a REPA-learned distribution may collapse onto fewer folds within its designable subset. Across most configurations, the designable subset is structurally more concentrated than the baseline’s, as measured by the mean pairwise TM-score (pwTM). We use different metrics for whole-set and designable-subset diversity to account for the different set sizes. fS-A is a distributional measure, but the designable subset is far smaller than the whole set, which makes it less interpretable in this context. pwTM is a pairwise average and more robust at smaller sizes. The concentration is especially severe for β -strand-rich samples (Table 6.6). Crucially, the pattern holds even at matched β -content, i.e. β -strand fraction. Within the most β -rich bin (β -strand fraction $\geq 25\%$), the baseline stays diverse (pwTM ≈ 0.13) while REPA variants collapse onto near-identical folds (≈ 0.87). Somehow, the REPA construction appears to generate the *same* β -rich folds repeatedly, while the baseline spans many.

On PDB, fold concentration occurs early in training, and the baseline grows out of it. After 700K training steps, the PDB baseline and its REPA variants alike concentrate their β -rich designable samples. However, by 1.0M steps, the baseline has diversified while REPA stays concentrated (Table 6.6). This suggests that REPA does not *engender* the concentration, but rather fails to grow out of the early-training regime that the baseline leaves behind.

Different encoders nudge the model into different distributions. Notably, the AFDB-MPNN variant does not display fold-concentration: its designable subset stays as diverse as the baseline’s in every β -bin (Table 6.6). One possible explanation here is that encoder routing leads to different secondary- and tertiary-structure distributions. This MPNN configuration appears to nudge the learned distribution towards helices, which may be why it escapes this largely strand-heavy phenomenon (Appendix A.7).

Metric (700K)	Sampler setting					
	ODE	$\gamma=0$	$\gamma=0.35$	$\gamma=0.45$	$\gamma=0.5$	$\gamma=1.0$
<i>PDB-trained — baseline vs. REPA MPNN-L4</i>						
FPSD ↓, base	353	782	462	437	425	508
FPSD ↓, REPA	279	561	371	349	350	424
Des ↑, base	0.04	0.84	0.56	0.46	0.36	0.01
Des ↑, REPA	0.16	0.78	0.66	0.64	0.52	0.11
pwTM (des.) ↓, base	0.26	0.75	0.29	0.27	0.24	0.43
pwTM (des.) ↓, REPA	0.33	0.44	0.36	0.35	0.34	0.34
<i>AFDB-trained — baseline vs. REPA GearNet-L4</i>						
FPSD ↓, base	220	641	522	494	475	366
FPSD ↓, REPA	148	586	328	298	292	276
Des ↑, base	0.12	0.91	0.72	0.72	0.69	0.21
Des ↑, REPA	0.35	0.96	0.84	0.77	0.72	0.28
pwTM (des.) ↓, base	0.20	0.69	0.25	0.22	0.22	0.15
pwTM (des.) ↓, REPA	0.19	0.68	0.31	0.26	0.25	0.17

Table 6.7: **REPA tracks the same designability–diversity trade-off as the baseline across sampler settings, at a higher level.** (Absolute values at step 700K; baseline vs. a representative REPA variant per dataset; pwTM is over the designable subset; green/red = REPA better/worse than baseline, small $|\Delta|$ uncoloured; $\gamma=0.45$ is Proteina’s default; $n \leq 256$.)

Nevertheless, REPA’s whole-set gains persist deep into training. At 1.0M steps, the strongest variant on each dataset keeps its fidelity and coverage lead — REPA-MPNN-L9 on PDB (FPSD -81 , fold-score $+0.29$) and REPA-GN-L4 on AFDB (-228 , $+3.68$) — even as their distinct designable-fold counts fall below baseline (-51 and -39 clusters).

Does alignment outlive its purpose? Given the early acceleration in generation quality induced by REPA, and the later-stage fold concentration we observe, it is possible that the best approach may be to stop alignment after early training. A similar arc is reported in the image diffusion literature, where REPA is found to help early and is best early-stopped [13]. We do not test such a schedule and make no argument of this form. However, on this reading, an early stop may let REPA diversify like the baseline while keeping the whole-set gains.

6.2.5 REPA’s gains hold across sampler noise levels

REPA navigates the same noise–quality trade-off as the baseline. As outlined in §2.1, Proteina samples with an SDE whose noise scale γ trades designability for diversity. This trade-off is an inherent property of the sampler, and REPA does not reshape it: as γ rises, REPA and the baseline both shed designability and gain diversity together (Table 6.7). What differs is the envelope each traces while doing so, as the next result shows.

REPA’s gains hold across the sampler’s noise range. We report our headline

Variant (700K)	ODE designability
<i>PDB-trained</i> (baseline = absolute; REPA = Δ vs. baseline)	
Baseline	0.04
REPA L4-random	-0.02
REPA L4-GN	+0.16
REPA L9-GN	+0.04
REPA L4-MPNN	+0.12
<i>AFDB-trained</i>	
Baseline	0.12
REPA L4-GN	+0.23
REPA L9-GN	+0.21
REPA L9-MPNN	+0.10

Table 6.8: **REPA’s learned distribution is more designable even without sampling noise.** The ODE limit samples the trunk’s learned distribution directly, which suggests REPA learns a better underlying distribution rather than merely benefiting from sampler noise. (ODE designability at step 700K; baseline absolute, REPA rows Δ ; green = higher, $|\Delta| < 0.025$ uncoloured. Single-seed.)

results at $\gamma=0.45$, the balance Proteina recommends. The advantage is not specific to that setting. For a representative variant on each dataset, REPA improves designability and fidelity at almost every γ (Table 6.7). The lone exception is the degenerate $\gamma=0$ regime, where the baseline saturates designability by repeating a few folds. REPA navigates the same trade-off as the baseline, but pushes the envelope outward: at almost every γ it reaches a higher level of both designability and FPSD than the baseline can.

REPA raises the deterministic-sampling floor, which points to a better learned distribution. The probability-flow ODE samples the trunk’s learned distribution without injected noise (§2.1). As such, it is the closest view of the true learned model distribution, before any correction. Under the baseline architecture, that distribution is barely designable: Des 0.04 on PDB and 0.12 on AFDB. The SDE sampling regime rescues this, with baseline designability climbing to ≈ 0.84 on PDB at $\gamma=0$. However, REPA’s learned distribution is already far more designable at ODE, across nearly every learned encoder (Table 6.8). Indeed, only the PDB random control fails to lift it (-0.02). This suggests REPA mechanistically improves the distribution the trunk learns.

6.2.6 What we do not claim

Several tempting claims do not survive our data, and we flag them here.

- *No asymptotic generation quality win on experimental data.* REPA accelerates PDB convergence, but the lead is transient. The baseline matches REPA’s FPSD plateau by ~ 1.4 M steps and edges below it by 1.6M (§6.2.3). The durable advantage on fidelity is on synthetic (AFDB) data, where the baseline plateaus poorly.
- *Unclear designability gains on synthetic data.* At the recommended sampler setting, AFDB designability is already biased upwards by construction and near-saturated

for the baseline, so a reliable whole-set gain is hard to establish (§6.1.1). The deterministic ODE floor of §6.2.5 is the exception.

- *No reliable novelty gain.* The signal is too noisy to call in either direction.
- *No general “ β -preservation”.* The secondary-structure shift is encoder $\times$ dataset conditional, and MPNN-AFDB reverses it (§6.2.4).
- *Co-movement, not mediation.* Representation and generation quality move together. We do not claim one formally causes the other, but we note that the pattern is consistent with the representation bottleneck hypothesis (§6.2.3).
- *Unconditional generation only.* Every experiment here is unconditional: we specify a target length but never a fold class or other structural motif. We do not claim these dynamics carry over to the conditional setting (§2.1).
- *No claim of optimality.* We do not search the configuration space for the best setup of layer, weight, or encoder. The question here is whether the mechanism works at all, not how to tune it for maximum gain (§6.1).

6.2.7 Robustness across scale

At smaller scale the learned-versus-random gap narrows, so we anchor headline claims at $n \leq 256$. At $n \leq 128$ the random control helps nearly as much on representation and whole-set metrics. The direction of every gain is consistent across scale, but the margin of learned over random shrinks, so we read $n \leq 128$ as a scale-robustness check rather than a headline regime. Training-dynamics ablations (λ , batch size, projector depth) are reported in Appendix A; length extrapolation — generating proteins longer than the training crop — is the natural next test and remains future work.

6.3 Discussion

This chapter establishes that REPA accelerates Proteina’s convergence on generation quality, and shows how the acceleration works. REPA is not a single lever, but a family of encoder-routed interventions. Across all variants we tested, models learn to represent structural characteristics that the flow-matching loss never learns in isolation. It also pulls the trunk’s geometry closer to the encoder’s, and this alignment generalises across other domain-specific encoders. Learned encoders consistently outperform our random control, a sign that the effect is genuinely mechanistic. Moreover, the representation and generation gains are encoder-matched, strongly suggesting the link is encoder-driven.

Designability and diversity form a known trade-off, and REPA pushes the envelope of that trade-off beyond where the baseline can reach by tuning the sampler alone. Sampler noise only slides a fixed model along its trade-off; REPA enhances the model itself, moving the whole frontier outward. The REPA architectures

we tested improve two things at once about the *whole* distribution learned over training: fidelity to reference and structural diversity. The clearest sign is the noise-free ODE floor, which REPA lifts, suggesting a genuinely better learned distribution. The gain carries an encoder-specific cost, as some encoders may concentrate the *designable* region of the learned distribution onto fewer folds.

REPA helps most where the baseline struggles most. On experimental data the acceleration is real but transient. The baseline is a strong asymptotic learner and eventually catches up. On synthetic data the baseline plateaus poorly, and there the lead is durable. The intervention is the same in both cases. The asymmetry tracks the baseline’s own convergence, not REPA. This is really a data-efficiency statement. It speaks to the motivation of this report directly. REPA’s payoff is largest in the abundant-but-synthetic regime the field increasingly leans on. We temper this on designability alone, where the synthetic-data proxy shares lineage with the training distribution (§6.1.1).

These results point past Proteina, to a thread we develop in the conclusion. In a domain whose representation is multi-factorial, choosing an alignment target is a routing decision, not a default. There is no single best encoder. Each one buys a different representational factor — a choice the sampler dial cannot make for you.

We carry this — and the contrast with the small-molecule study of Chapter 5, where the same intervention left no mark — into Chapter 7.

Chapter 7

Conclusions

Bibliography

- [1] Jonathan Ho, Ajay Jain, Pieter Abbeel. *Denoising Diffusion Probabilistic Models*. NeurIPS 2020. arXiv:2006.11239.
- [2] Yaron Lipman, Ricky T. Q. Chen, Heli Ben-Hamu, Maximilian Nickel, Matt Le. *Flow Matching for Generative Modeling*. ICLR 2023. arXiv:2210.02747.
- [3] Peter Holderrieth and Ezra Erives. *An Introduction to Flow Matching and Diffusion Models*. MIT 6.S184 lecture notes, 2025. diffusion.csail.mit.edu.
- [4] Yaron Lipman, Marton Havasi, Peter Holderrieth, Neta Shaul, Matt Le, Brian Karrer, Ricky T. Q. Chen, David Lopez-Paz, Heli Ben-Hamu, Itai Gat. *Flow Matching Guide and Code*. 2024. arXiv:2412.06264.
- [5] Sihyun Yu, Sangkyung Kwak, Huiwon Jang, Jongheon Jeong, Jonathan Huang, Jinwoo Shin, Saining Xie. *Representation Alignment for Generation: Training Diffusion Transformers Is Easier Than You Think*. ICLR 2025. arXiv:2410.06940.
- [6] Jaskirat Singh, Xingjian Leng, Zongze Wu, Liang Zheng, Richard Zhang, Eli Shechtman, Saining Xie. *What matters for Representation Alignment: Global Information or Spatial Structure?* 2025. arXiv:2512.10794.
- [7] Chenyu Wang, Cai Zhou, Sharut Gupta, Zongyu Lin, Stefanie Jegelka, Stephen Bates, Tommi Jaakkola. *Learning Diffusion Models with Flexible Representation Guidance*. 2025. arXiv:2507.08980.
- [8] Hyosoon Jang, Hyunjin Seo, Yunhui Jang, Seonghyun Park, Sungsoo Ahn. *Boltz is a Strong Baseline for Atom-level Representation Learning*. 2026. arXiv:2602.13249.
- [9] Xiangdong Zhang, Jiaqi Liao, Shaofeng Zhang, Fanqing Meng, Xiangpeng Wan, Junchi Yan, Yu Cheng. *VideoREPA: Learning Physics for Video Generation through Relational Alignment with Foundation Models*. 2025. arXiv:2505.23656.
- [10] Tri Ton, Jiajing Hong, Chang D. Yoo. *Temporally Aligned Audio for Video with Autoregression*. ICCV 2025. arXiv:2504.05684.
- [11] Wei Zhang, Shijie Jin, Hangrui Zhang, Yi Li, Yan Wang. *CrystalREPA: Element-aware Representation Alignment for Crystal Generation*. 2026. arXiv:2605.08960.
- [12] Runqian Wang and Kaiming He. *Diffuse and Disperse: Image Generation with Representation Regularization*. 2025. arXiv:2506.09027.

- [13] Ziqiao Wang, Wangbo Zhao, Yuhao Zhou, Zekai Li, Zhiyuan Liang, Mingjia Shi, Xuanlei Zhao, Pengfei Zhou, Kaipeng Zhang, Zhangyang Wang, Kai Wang, Yang You. *REPA Works Until It Doesn't: Early-Stopped, Holistic Alignment Supercharges Diffusion Training*. 2025. arXiv:2505.16792.
- [14] Tomas Geffner et al. *Proteina: Scaling Flow-based Protein Structure Generative Models*. ICLR 2025. arXiv:2503.00710.
- [15] C. Vonessen et al. *TABASCO: A Fast, Simplified Model for Molecular Generation with Improved Physical Quality*. 2025. arXiv:2507.00899.
- [16] Marcus Olivecrona, Thomas Blaschke, Ola Engkvist, Hongming Chen. "Molecular de-novo design through deep reinforcement learning". *Journal of Cheminformatics* 9:48 (2017). doi:10.1186/s13321-017-0235-x.
- [17] Alex Zhavoronkov et al. "Deep learning enables rapid identification of potent DDR1 kinase inhibitors". *Nature Biotechnology* 37 (2019), pp. 1038–1040. doi:10.1038/s41587-019-0224-x.
- [18] Joseph L. Watson et al. "De novo design of protein structure and function with RFDiffusion". *Nature* 620 (2023), pp. 1089–1100. doi:10.1038/s41586-023-06415-8.
- [19] John B. Ingraham et al. "Illuminating protein space with a programmable generative model". *Nature* 623 (2023), pp. 1070–1078. doi:10.1038/s41586-023-06728-8.
- [20] Amil Merchant et al. "Scaling deep learning for materials discovery". *Nature* 624 (2023), pp. 80–85. doi:10.1038/s41586-023-06735-9.
- [21] Claudio Zeni et al. "A generative model for inorganic materials design". *Nature* 639 (2025), pp. 624–632. doi:10.1038/s41586-025-08628-5.
- [22] Liang Wang, Chao Song, Zhiyuan Liu, Yu Rong, Qiang Liu, Shu Wu, Liang Wang. *Diffusion Models for Molecules: A Survey of Methods and Tasks*. 2025. arXiv:2502.09511.
- [23] Zeming Lin et al. "Evolutionary-scale prediction of atomic-level protein structure with a language model". *Science* 379.6637 (2023), pp. 1123–1130. doi:10.1126/science.ade2574.
- [24] J. Dauparas et al. "Robust deep learning-based protein sequence design using ProteinMPNN". *Science* 378.6615 (2022), pp. 49–56. doi:10.1126/science.add2187.
- [25] Zuobai Zhang, Minghao Xu, Arian Jamasb, Vijil Chenthamarakshan, Aurelie Lozano, Payel Das, Jian Tang. *Protein Representation Learning by Geometric Structure Pre-training*. ICLR 2023. arXiv:2203.06125.
- [26] Arian R. Jamasb, Alex Morehead, Chaitanya K. Joshi, et al. *Evaluating Representation Learning on the Protein Structure Universe*. ICLR 2024.
- [27] John Ingraham, Vikas K. Garg, Regina Barzilay, Tommi Jaakkola. "Generative Models for Graph-Based Protein Design". *Advances in Neural Information Processing Systems (NeurIPS)* 32 (2019).

- [28] Martin Steinegger, Johannes Söding. “MMseqs2 enables sensitive protein sequence searching for the analysis of massive data sets”. *Nature Biotechnology* 35.11 (2017), pp. 1026–1028. doi:10.1038/nbt.3988.
- [29] Ilyes Batatia et al. “A foundation model for atomistic materials chemistry”. *The Journal of Chemical Physics* 163.18 (2024), 184110. doi:10.1063/5.0155322.
- [30] Jackson W. Burns et al. “Deep Learning Foundation Models from Classical Molecular Descriptors”. arXiv preprint (2025). arXiv:2506.15792.
- [31] Quentin Garrido, Randall Balestrieri, Laurent Najman, Yann LeCun. *RankMe: Assessing the Downstream Performance of Pretrained Self-Supervised Representations by Their Rank*. ICML 2023. arXiv:2210.02885.
- [32] Michael Heinzinger et al. “Bilingual language model for protein sequence and structure”. *NAR Genomics and Bioinformatics* 6.4 (2024), lqae150. doi:10.1093/nargab/lqae150.
- [33] Jin Su et al. “SaProt: Protein Language Modeling with Structure-aware Vocabulary”. *bioRxiv* (2023). Preprint. doi:10.1101/2023.10.01.560349.
- [34] Christoph Schuhmann et al. *LAION-5B: An open large-scale dataset for training next generation image-text models*. NeurIPS 2022 Datasets and Benchmarks. arXiv:2210.08402.
- [35] Helen M. Berman et al. “The Protein Data Bank”. *Nucleic Acids Research* 28.1 (2000), pp. 235–242. Statistics retrieved Nov 2025, rcsb.org/stats.
- [36] Mihaly Varadi et al. “AlphaFold Protein Structure Database in 2024: providing structure coverage for over 214 million protein sequences”. *Nucleic Acids Research* 52 (2024), D368–D375. doi:10.1093/nar/gkad1011.
- [37] Simon Axelrod and Rafael Gómez-Bombarelli. “GEOM, energy-annotated molecular conformations for property prediction and molecular generation”. *Scientific Data* 9, 185 (2022). doi:10.1038/s41597-022-01288-4.
- [38] Lorna Richardson et al. “MGnify: the microbiome sequence data analysis resource in 2023”. *Nucleic Acids Research* 51 (2023), D753–D759. doi:10.1093/nar/gkac1080. Release statistics retrieved May 2022 (2.4 billion non-redundant sequences): ebi.ac.uk.
- [39] Luciano Brocchieri and Samuel Karlin. “Protein length in eukaryotic and prokaryotic proteomes”. *Nucleic Acids Research* 33.10 (2005), pp. 3390–3400. doi:10.1093/nar/gki615.
- [40] Michael M. Bronstein, Joan Bruna, Taco Cohen, Petar Veličković. *Geometric Deep Learning: Grids, Groups, Graphs, Geodesics, and Gauges*. 2021. arXiv:2104.13478.
- [41] Alexandre Duval, Simon V. Mathis, Chaitanya K. Joshi, Victor Schmidt, Santiago Miret, Fragkiskos D. Malliaros, Taco Cohen, Pietro Liò, Yoshua Bengio, Michael Bronstein. *A Hitchhiker’s Guide to Geometric GNNs for 3D Atomic Systems*. 2024. arXiv:2312.07511.

- [42] Nathaniel Thomas, Tess Smidt, Steven Kearnes, Lusann Yang, Li Li, Kai Kohlhoff, Patrick Riley. *Tensor Field Networks: Rotation- and Translation-Equivariant Neural Networks for 3D Point Clouds*. 2018. arXiv:1802.08219.
- [43] Alexey Dosovitskiy et al. *An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale*. ICLR 2021. arXiv:2010.11929.
- [44] Jinwoo Kim, Tien Dat Nguyen, Seonwoo Min, Sungjun Cho, Moontae Lee, Honglak Lee, Seunghoon Hong. *Pure Transformers are Powerful Graph Learners*. NeurIPS 2022. arXiv:2207.02505.
- [45] John Jumper et al. “Highly accurate protein structure prediction with AlphaFold”. *Nature* 596 (2021), pp. 583–589. doi:10.1038/s41586-021-03819-2.
- [46] Josh Abramson et al. “Accurate structure prediction of biomolecular interactions with AlphaFold 3”. *Nature* 630 (2024), pp. 493–500. doi:10.1038/s41586-024-07487-w.
- [47] Johann Brehmer et al. *Does equivariance matter at scale?* NeurIPS 2024. arXiv:2410.23179.
- [48] Jason Yim et al. *Fast protein backbone generation with SE(3) flow matching*. 2023. arXiv:2310.05297.
- [49] Yeqing Lin and Mohammed AlQuraishi. *Generating Novel, Designable, and Diverse Protein Structures by Equivariantly Diffusing Oriented Residue Clouds*. ICML 2023. arXiv:2301.12485.
- [50] Tuan Le et al. *Navigating the design space of equivariant diffusion-based generative models for de novo 3D molecule generation*. ICLR 2024. arXiv:2309.17296.
- [51] Clement Vignac et al. *MiDi: Mixed Graph and 3D Denoising Diffusion for Molecule Generation*. ECML PKDD 2023. arXiv:2302.09048.
- [52] Ross Irwin et al. *Efficient 3D Molecular Generation with Flow Matching and Scale Optimal Transport*. 2024. arXiv:2406.07266.
- [53] National Institutes of Health. *Regulatory Knowledge Guide for Small Molecules*. NIH SEED, 2024. seed.nih.gov.
- [54] Encyclopaedia Britannica. *What is the difference between a peptide and a protein?* britannica.com.
- [55] Gisbert Schneider. “Automating drug discovery”. *Nature Reviews Drug Discovery* 17 (2018), pp. 97–113. doi:10.1038/nrd.2017.232.
- [56] Martin Buttenschoen, Garrett M. Morris, and Charlotte M. Deane. “PoseBusters: AI-based docking methods fail to generate physically valid poses or generalise to novel sequences”. *Chemical Science* 15 (2024), pp. 3130–3139. arXiv:2308.05777.
- [57] Kristina Preuer, Philipp Renz, Thomas Unterthiner, Sepp Hochreiter, and Günter Klambauer. “Fréchet ChemNet Distance: A Metric for Generative Models for

- Molecules in Drug Discovery”. *Journal of Chemical Information and Modeling* 58.9 (2018), pp. 1736–1741. doi:10.1021/acs.jcim.8b00234.
- [58] Martin Heusel, Hubert Ramsauer, Thomas Unterthiner, Bernhard Nessler, and Sepp Hochreiter. “GANs Trained by a Two Time-Scale Update Rule Converge to a Local Nash Equilibrium”. *Advances in Neural Information Processing Systems (NeurIPS)* 30 (2017). arXiv:1706.08500.
- [59] G. Richard Bickerton et al. “Quantifying the chemical beauty of drugs”. *Nature Chemistry* 4 (2012), pp. 90–98. doi:10.1038/nchem.1243.
- [60] Greg Landrum et al. *RDKit: Open-source cheminformatics*. www.rdkit.org.
- [61] Pascal Notin et al. “Machine learning for functional protein design”. *Nature Biotechnology* 42 (2024), pp. 216–228. doi:10.1038/s41587-024-02127-0.
- [62] Yeqing Lin, Minji Lee, Zhao Zhang, and Mohammed AlQuraishi. “Out of Many, One: Designing and Scaffolding Proteins at the Scale of the Structural Universe with Genie 2”. *arXiv preprint* arXiv:2405.15489 (2024). arXiv:2405.15489.
- [63] Christine A. Orengo, A. D. Michie, S. Jones, David T. Jones, M. B. Swindells, Janet M. Thornton. “CATH — a hierarchic classification of protein domain structures”. *Structure* 5.8 (1997), pp. 1093–1108. doi:10.1016/S0969-2126(97)00260-8.
- [64] Minyoung Huh, Brian Cheung, Tongzhou Wang, Phillip Isola. *The Platonic Representation Hypothesis*. ICML 2024. arXiv:2405.07987.

Appendix A

Technical details

A.1 Dataset composition and split overlap

Table A.1 summarises the two datasets used in Chapter 6. Both training sets are α/β -dominated and span a substantial fraction of CATH architectures (44.2% PDB, 61.8% AFDB), which is what makes the per-chain CATH fold probe meaningful. Neither random split is sequence-clustered: 98.3% of PDB and 92.2% of AFDB validation chains have a $\geq 30\%$ -identity homolog in their own training set, so absolute in-house scores are inflated and we read rank-order and Δ -from-baseline instead. The cross-database overlap is far lower — only 62.1% of AFDB validation chains have a $\geq 30\%$ -identity hit in PDB training — so a homology-filtered AFDB set serves as a cleaner probe set for PDB-trained models (Chapter 3).

A.2 Choice of CATH fold-probe level

We report CATH *architecture* (CATH-A) as the headline fold-probe level throughout Chapter 6. Table A.2 justifies that choice: Class is too coarse (four buckets, one already near-saturated), Topology too finely-bucketed to probe reliably, and Architecture sits between — discriminative but learnable.

Property	PDB	AFDB (Swiss-Prot)
Origin	Experimental (X-ray, cryo-EM)	Predicted (AlphaFold2, v6)
Residue-length crop n	≤ 256 residues	
Train / val / test chains	273,771 / 3,190 / 323	229,670 / 4,521 / 235
CATH coverage (train)	44.2%	61.8%
Val $\leftrightarrow$ train, $\geq 30\%$ id.	98.3%	92.2%
AFDB-val $\leftrightarrow$ PDB-train, $\geq 30\%$ id.	62.1%	

Table A.1: Properties of the two datasets we study, at the headline $n \leq 256$ regime. Sequence identity is computed with MMseqs2 [28] (**easy-search**, $\geq 30\%$ identity over $\geq 80\%$ alignment coverage).

CATH level	#classes	baseline probe acc.	verdict
C (class)	4	0.733	one class (α/β) is 55% of the data; coarse, already near-saturated — <i>gameable</i>
A (architecture)	~ 40	0.407	balanced and discriminative, with headroom to move — our headline fold metric
T (topology)	~ 1400	0.301	long-tailed (top fold $\sim 18\%$; only ~ 89 classes populated enough to probe); per-class accuracy is unreliable — <i>too many buckets</i>

Table A.2: **We read fold-level representation quality off CATH *architecture* — Class is too coarse to discriminate, Topology too finely-bucketed to probe reliably.** Class has 4 buckets (one at 55%, baseline already 0.73); Topology has ~ 1400 imbalanced buckets; Architecture sits between — enough classes to be meaningful, enough examples per class to be learnable — so CATH-A is our headline fold metric (C and T are reported in the appendix). (“Baseline probe acc.” is a linear probe on the un-aligned trunk.)

A.3 Leakage-controlled probe evaluation

The random train/val split (§A.1) is not sequence-clustered, so probe scores are exposed to two distinct leakages. *Model leakage* is asymmetric: a model trained on the leaky split has effectively seen near-duplicates of the evaluation proteins, an advantage a differently-curated model (such as an externally-pretrained checkpoint) does not share. *Probe leakage* is symmetric: a probe fit on training chains near-identical to its evaluation set can shortcut by matching sequences rather than learning generalisable structure, inflating every model’s absolute score.

We control these by filtering the evaluation. A *probe-clean* split removes probe leakage alone — the probe’s training chains are filtered to have no $\geq 30\%$ sequence-identity hit to the evaluation set, which is otherwise the full validation set. A *blinded* split removes both — the evaluation proteins are held to under 30% identity to all training data (here, a cross-database set of AFDB chains with no such hit in either the PDB or AFDB training pool, ~ 325 proteins). We use the blinded split for the per-residue probes (inverse folding, dihedral), where the ~ 325 proteins still yield $\sim 43\text{k}$ residues — full statistical power. The per-chain CATH probe is different: fold classification needs many distinct chains spread across populated fold classes, and the blinded slice is too small and too class-sparse for this, so we relax it to the probe-clean split on the fuller validation set ($\sim 3,190$ chains). Comparing the splits also decomposes the leakage — the dirty-to-probe-clean drop isolates probe leakage and the probe-clean-to-blinded drop isolates model leakage — and the REPA-over-baseline gap survives in every regime.

	GearNet	ProteinMPNN	ESM2
Baseline	0.6 ^{L6}	0.9 ^{L7}	2.7 ^{L8}
REPA GN-L4	2.6^{L4†}	2.9 ^{L6}	8.2 ^{L8}
REPA GN-L9	22.9^{L8†}	26.7 ^{L8}	40.1 ^{L8}
REPA MPNN-L4	4.8 ^{L7}	16.6^{L7†}	19.8 ^{L5}
REPA MPNN-L9	5.0 ^{L7}	14.2^{L6†}	14.2 ^{L7}

Table A.3: **Every REPA variant raises per-residue alignment to all three encoders, including the two it was never aligned to.** Across-layer peak per-residue CKNNA (bootstrap median, $\times 10^3$), with the peak trunk layer as a superscript; $k=10$, step 1.0M, $n \leq 256$ PDB-trained. The † marks each model’s own alignment target. Every REPA row exceeds the baseline in every column, the Platonic-convergence signature discussed in §6.2.2. Absolute values are small (an order of magnitude below image-domain REPA), so we read the pattern, not the magnitude.

A.4 Representation-alignment matrix

The alignment diagnostic in §6.2.2 reports peak per-residue CKNNA for a single illustrative variant. Table A.3 gives the full model $\times$ encoder matrix behind that reading. It confirms the off-diagonal generalisation we describe there. Aligning the trunk to one encoder raises its CKNNA to the other two as well, not just the alignment target. The peak sits at the last trunk layers (L7–L8 here). The absolute values stay an order of magnitude below image-domain REPA, so we read the direction and layer pattern rather than the raw magnitude.

A.5 Representation–generation coupling on AFDB

Section 6.2.3 establishes the representation–generation coupling on PDB-trained models (Table 6.4). Table A.4 repeats the analysis for AFDB-trained models. The coupling holds here too, and in the same encoder-matched form. Local probes (inverse folding, dihedral) track designability, while the fold probe (CATH-A) tracks FPSD. The off-axis pairs are weak or absent, the signature of an encoder-routed effect rather than a generic one. We read AFDB designability with the proxy caveat of §6.1.1, so we lean on the FPSD column and the routing pattern rather than the absolute designability correlation.

A.6 The ssJSD-2D secondary-structure metric

Proteina’s secondary-structure metric compares the *mean* helix and strand content of a generated set against a reference. This collapses each set to a single point, so it cannot see variance, bimodality, or tails — a model that matches the average composition while collapsing its spread scores the same as one that reproduces it. ssJSD-2D removes this blind spot by comparing the full joint distribution.

For each backbone we annotate every residue as helix, strand, or coil with a C α -only

Representation quality	FPSD		Designability	
	partial	raw	partial	raw
Inverse folding (top-1)	-0.17	-0.16	+0.50	+0.50
Dihedral MAE	+0.03	-0.03	+0.37	+0.09
CATH-A	+0.42	+0.43	+0.29	+0.34

Table A.4: **On AFDB too, better trunk representations track better generation.** The AFDB counterpart to Table 6.4: partial Pearson controlling for training step (raw alongside), oriented so positive means better representation accompanies better generation. Generation is FPSD-AFDB and designability; representation probes are read at $n_{\text{train}}=1000$ on the cross-database blinded set. ($n=40$ AFDB checkpoints pooled over the baseline, the REPA variants, and the random control.)

secondary-structure assignment (P-SEA, as implemented in Biotite) and record the fraction of residues in helix (f_H) and in strand (f_E). A set of N backbones is then a cloud of N points in the (f_H, f_E) plane. We bin that plane into a 5×5 grid over $[0, 1]^2$, form normalised two-dimensional histograms for the generated and reference sets, and report the Jensen–Shannon divergence between them (lower is better). Binning the joint distribution, rather than comparing marginal means, lets the metric reward a model for reproducing the *shape* of the secondary-structure distribution and penalise mode collapse onto, say, all- α backbones.

A.7 Secondary-structure composition of the designable subset

The encoder-routing reading in §6.2.4 — GearNet shifting the designable subset toward β -strand content, MPNN on AFDB away from it — rests on the secondary-structure composition of the designable backbones (Table A.5). The shifts are small in absolute terms but seed-consistent: on AFDB the GearNet-L9 and MPNN-L9 strand and helix ranges do not overlap across seeds, in opposite directions. The main-text E%des figures (Table 6.5) are the per-seed means of the strand column; we report the spread here rather than in the main table to avoid clutter.

A.8 Full encoder profiles

Tables 4.2 and 4.3 in the main text show only the diagnostics that drive each verdict. Tables A.6 and A.7 give the complete set of measured values for every candidate encoder, including the columns omitted from the main text (atom/context probes, participation ratio, and the best-projector cosine the lift is computed from).

Variant (700K, $\gamma=0.45$)	Strand f_E	Helix f_H	Seeds
<i>PDB-trained</i>			
Baseline	0.22 ± 0.02	0.48 ± 0.03	3
REPA L4-random	0.11 ± 0.01	0.61 ± 0.02	3
REPA L4-GN	0.22 ± 0.01	0.40 ± 0.02	3
REPA L9-GN	0.18 ± 0.02	0.48 ± 0.03	3
REPA L4-MPNN	0.23 ± 0.01	0.39 ± 0.03	3
REPA L9-MPNN	0.15 ± 0.01	0.54 ± 0.02	3
<i>AFDB-trained</i>			
Baseline	0.16 ± 0.01	0.47 ± 0.01	2
REPA L4-GN	0.14 ± 0.00	0.47 ± 0.01	2
REPA L9-GN	0.19 ± 0.01	0.38 ± 0.03	3
REPA L9-MPNN	0.11 ± 0.00	0.55 ± 0.01	2

Table A.5: **Secondary-structure composition of the designable subset (step 700K, $\gamma=0.45$).** Mean strand (f_E) and helix (f_H) fractions over the designable backbones, $\pm$ half the seed range. GearNet shifts the designable subset toward strand and away from helix; MPNN on AFDB does the reverse. The AFDB GearNet-L9 vs. MPNN-L9 ranges are non-overlapping on both axes. The strand column is the absolute form of the E%des column in Table 6.5, whose entries are the per-seed means reported here. (AFDB MPNN-L4 omitted: no multi-seed $\gamma=0.45$ data.)

Encoder	Dim	atom probe	context- Δ	3D-aware	Floor	Lift
CheMeleon (2D graph)	2048	1.000	0.093	no	0.43	+0.04
MACE (3D potential)	192	1.000	0.260	weak	0.86	+0.005

Encoder	Dim	Sparsity	RankMe
CheMeleon (2D graph)	2048	93.8%	1166
MACE (3D potential)	192	0.0%	40.6

Table A.6: Full molecule encoder profile. *atom probe*: linear-probe accuracy for atom type (trivial in QM9, so omitted from the main table). *context- Δ* : same-element, different-environment cosine spread. *Floor/Lift* as in Table 4.2. RankMe is the Roy–Vetterli effective rank (GEOM value).

Encoder	Dim	AA-probe	context- Δ	Floor	Best	Lift	RankMe	PR
GearNet (CATH-trained)	512	0.137	0.222	0.425	0.434	+0.009	256.4	37.9
<i>random-init</i>	512	0.127	0.035	0.952	0.954	+0.003	99.8	1.5
ProteinMPNN (inv-fold)	128	0.327	0.022	0.835	0.850	+0.014	84.9	33.3
ESM2-650M (last layer)	1280	0.998	0.098	0.671	0.723	+0.053	1030.5	43.3
PW-GearNet (torsional)	3072	0.928	0.102	0.710	0.723	+0.013	207.7	5.5
MC-GearNet-Edge	3072	0.146	0.043	0.855	0.858	+0.002	12.0	1.0

Table A.7: Full protein encoder profile. *AA-probe*, *context- Δ* (within-minus-between protein cosine), *Floor* (mean-direction cosine), *Best* (best-projector cosine), *Lift* (= Best – Floor), *RankMe* and *PR* (participation ratio). MC-GearNet’s mean L2 norm is 1.5×10^6 , with 95% of its variance in a single dimension.