

n=256 convergence — generation metrics vs training step
Higher is better unless axis inverted; baseline drawn first.

