

n=256 convergence — generation metrics vs training step
Direction (higher vs lower better) varies by panel — see ylabel; baseline drawn first.

