

n=256 — Pareto frontiers (point size $\propto$ training step; thick line per method)
Top-right wins; REPA frontiers above-right of baseline = Pareto improvement.

