

n=256 — generation vs representation envelope (point size $\propto$ training step)
Top-right = better on both axes; lines connect same-run checkpoints in step order.

