

n=256 convergence — representation probe quality vs training step
Best layer per checkpoint (max for ↑ metrics, min for ↓); t=1.0; baseline drawn first.

